

A General Framework for Compositional Network Modeling

Ryan Beckett
Microsoft Research

Ratul Mahajan
University of Washington, Intentionet

ABSTRACT

We advocate for an approach to network modeling and analysis based on a common intermediate language. Unlike today, where each tool builds a custom model and analysis engine for its target network functionality, we argue that network functionality should be expressed in a common language. This approach makes it easier to expand formal analysis to new functionality and analyze interactions between dependent functionalities (e.g., routing and packet filtering). We demonstrate the feasibility of this approach by developing an intermediate language called Zen and three different analyses for programs in that language. For representative data plane and control plane functionalities, we find that Zen reduces the modeling effort by an order of magnitude, while providing analysis performance that matches custom tools.

CCS CONCEPTS

• **Networks** → *Network reliability*; • **Theory of computation** → *Program verification*; • **Software and its engineering** → *Model checking*; *Functional languages*.

KEYWORDS

network verification, intermediate verification language

ACM Reference Format:

Ryan Beckett and Ratul Mahajan. 2020. A General Framework for Compositional Network Modeling. In *Proceedings of the 19th ACM Workshop on Hot Topics in Networks (HotNets '20)*, November 4–6, 2020, Virtual Event, USA. ACM, New York, NY, USA, 8 pages. <https://doi.org/10.1145/3422604.3425930>

1 INTRODUCTION

As more devices and services connect to the Internet, computer networks continue to increase in scale and complexity. Formal modeling and analysis of network behavior has emerged as a key approach to managing such complexity, by finding problems before they cause an outage.

Researchers have developed systems and models to verify a range of network functionalities such as packet filtering (access control lists) [31], packet forwarding [22], routing [11], network address translation, and other types of packet transformations [42]. This development is also not limited to academia. Large cloud providers, such as Alibaba [36], Amazon [2], and Microsoft [20], are developing and deploying network verification systems.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

HotNets '20, November 4–6, 2020, Virtual Event, USA

© 2020 Association for Computing Machinery.

ACM ISBN 978-1-4503-8145-1/20/11...\$15.00

<https://doi.org/10.1145/3422604.3425930>




Figure 1: Network verification today. Each tool has its own model of target functionality and its own analysis approach.




Figure 2: Network verification with model composition.

However, as shown in Figure 1, each such tool today is a monolith, with its own model of the target functionality and its own analysis engine. The engine typically includes a front end that encodes the domain and a back-end solver that may be custom or standard solvers such as an SMT solver or BDDs (binary decision diagram). Any sharing between tools that exists is surface-level. For instance, Minesweeper [3] and ARC [12] use Batfish [11] to obtain a structured representation of network configuration files, but then build their own models and their own analysis engines. If Batfish is later extended to new functionality, these tools will not support that functionality unless they are updated as well.

This state of affairs makes it hard to expand verification to new network functionality because that expansion requires developing the full pipeline from scratch, a significant endeavor. Consequently, there is a substantial amount of network functionality that is not covered by any tool today, from the link layer (e.g., multiple access protocols) to the network layer (e.g., EIGRP routing protocol) to the application layer (e.g., HTTP firewalls and URL-based forwarding). The extent of unverified network functionality will only increase as cloud network providers continue to roll out new features and as engineers rapidly innovate atop programmable NICs and switches.

Making matters worse, it is not enough to have *some* tool to verify each piece of network functionality in isolation because the ultimate network behavior depends on the interactions of these pieces. When an individual piece is verified, it assumes that the pieces it depends upon are correct. If tools use disparate models, such assumptions can go unchecked, and bugs can lurk at the boundary of independently verified pieces.

We call for a compositional approach to network modeling and analysis. Illustrated in Figure 2, here, network functionality is modeled in a common language, and analysis engines target the language instead of a specific network functionality. This approach enables rapid expansion of verification to new functionality by separating concerns. Users need only encode domain-specific functionality in the modeling language, while authors of analysis engines need only target the modeling language and not a particular domain. It also enables holistic analysis across functional pieces. One can combine the models of multiple pieces to obtain a model of the joint behavior that can then help uncover bugs at the boundaries. These capabilities will put us on the path to fully verified networks, where all critical functionalities and interactions can be verified.

Our approach is inspired by the software analysis domain. There is a vibrant ecosystem around intermediate languages such as LLVM [25] and Boogie [27] that can encode the semantics of programs in multiple other source languages. A variety of analysis and optimization tools are available for these languages, which benefit all source languages. Boogie also became the basis for writing provably correct programs [26] and full-system verification [18] where each instruction was formally verified. Our intent is to similarly accelerate innovation for network modeling and analysis.

The success of our proposal depends on designing the right intermediate modeling language. It must be expressive enough to encode diverse and complex network functionality, restricted enough to permit efficient automated analysis, and simple enough to reduce the burden of implementing analyses.

We present the preliminary design of such an expressive yet compact modeling language called Zen. At its core, Zen is an expression-oriented language with basic types including booleans, integers, tuples, objects/structs, lists, and maps. To make it easier to both encode network functionality and author analysis tools, we embed Zen in the C# language. Users write functions that process Zen objects as they would write any C# code. We then use reflection to automatically analyze and translate objects into logic or a novel *state set* abstraction that we have developed.

We built several analysis engines for Zen including a simulator, bounded model checker, unbounded model checker, and test input generator. We further encoded the functionality of a number of network components such as route policies, IP GRE tunnels, and ACLs, and compare their expressiveness and performance with state-of-the-art tools. We find that Zen can often implement complex functionality in an order of magnitude less code and that its analysis is efficient.

2 MOTIVATION

To realize our vision of rapid development of network analyses, we need an intermediate verification language (IVL) that is: (1) compositional, and (2) general.

Objective #1: Compositional. Consider a setting where multiple virtual networks, called *overlays*, run atop a shared physical network, called the *underlay*. Such virtualized networks are the norm in modern data centers because they provide isolation and easy migration for overlays. Figure 3 shows one such network. V_a and V_b (virtual machines or containers) are overlay endpoints with a (virtual) link between them, and the underlay has three nodes.




Figure 3: An example virtualized network with an illustration of how packets flow across it.

There are many ways to virtualize a network, but independent of the technique used, overlay packets tend to be tunneled. So, as shown in the figure, when V_a sends a packet to V_b , this packet is encapsulated by U_1 within another header with U_3 as the destination. Then, the packet reaches U_3 via the underlay, and it is decapsulated and passed to V_b . The overlay and underlay have their own control and data plane systems. That is, the overlay has its own routing, forwarding, and packet filtering rules, and the underlay has its own version, though the underlay processing may be based on overlay headers as well. The two systems may be completely different, e.g., the overlay may use an SDN-style control plane and the underlay may use distributed routing protocols.

When verifying a virtualized network, it is desirable to verify the combined impact of overlay and underlay processing. Using today's approach, of building monolithic analysis tools that do not decouple network behavior modeling from solvers, a virtualized network may be verified using one of two ways. The first is to separately verify the overlay and the underlay, using tools appropriate for each network type. Here, overlay verification must assume that the underlay provides perfect connectivity, and the underlay verification will be agnostic to overlays that run atop it. The second method is to build and validate the a combined model of the overlay and underlay.

Both methods are problematic. The first cannot find problems that manifest when the two networks interact. For instance, the underlay may have a buggy packet filter that drops some types of overlay packets. This bug will not be found if we verify the underlay and the overlay separately. The second method has high engineering complexity, as it will have to model multiple types of overlay and underlay combinations.

Both methods are also hard to evolve because the solvers are intimately tied to the network model. Assume that a user introduces new functionality into the overlay (which is easy because developers can roll their own using a software update) or the underlay (which is also easy given the advent of programmable NICs, switches, and P4). Now, we will have to painstakingly update the models as well as the solvers.

Contrast the current approach with a compositional modeling approach. Here, the overlay and the underlay will be modeled separately. Once that is done, a range of solvers become available to verify the underlay and the overlay. It will also enable the creation of different types of combined models, which can then be verified using any available solver.

Our approach will find bugs at the intersection of the overlay and the underlay and will find overlay-only and underlay-only bugs faster. It will also be easier to evolve. Supporting new overlay

Analysis	Rosette	Kaplan	Boogie	NV	Zen
HSA [22]	✗	✗	✗	✓	✓
AP [41, 42]	✗	✗	✗	✗	✓
Anteater [29]	✓	✓	✓	✗	✓
Minesweeper [3]	✓	✓	✓	✓	✓
Bonsai [4]	✗	✗	✗	✗	✓
Shapeshifter [5]	✗	✗	✗	✓	✓

Table 1: Whether different intermediate verification languages (IVLs) can express example network analyses.

functionality only requires changing the overlay model, and the rest of the system stays as is.

While we considered virtualized networks above, the essential characteristics we discussed are universal. The full network behavior is a result of the interaction between multiple pieces (e.g., BGP plus OSPF plus access control lists), different networks have different combinations of these pieces, and existing pieces evolve and new pieces appear constantly.

Objective #2: General. Our IVL must be able to express a wide range of network analyses. This is where other IVLs for general-purpose software such as Kaplan, Boogie, and Rosette [24, 27, 37] fall short. They cannot express many verification analyses that have been highly effective for networks. A key reason is that network verification analyses often manipulate *sets* of objects (e.g., packets) whereas these IVLs work via compilation to logical constraints (which help find a counter example). There is a recent IVL for networks, called NV [13]. It too cannot express a wide range of analyses. Instead, NV bundles a few types of analyses, and modifications are needed to support more types. Table 1 shows examples of network analyses and whether different IVLs can support them. Zen achieves generality via *i*) a new *state set* abstraction for reasoning about sets of objects; and *ii*) its embedding as a library in a general purpose language, which allows users to express analyses by manipulating sets of objects using arbitrarily complex code.

3 MODELING NETWORKS WITH ZEN

We now describe how to model networks in Zen using the virtual network example above and then demonstrate how to analyze such models in the next section.

Encoding a domain. Assume that we want to build a data plane model for a network with longest-prefix match based forwarding and access control (ACLs) in both the overlay and underlay, and IP GRE [17] for tunneling. This task can be split into two parts. We first create ordinary classes in the host language (C#) that model objects such as packets, forwarding tables, and GRE tunnels. Figure 4 shows examples for IPv4 packets (line 1) and for packets with two headers, an overlay header and an optional underlay header (line 9).

After defining the objects to model, we encode the domain semantics by writing functions that process these objects. For example, to encode packet forwarding, one might write the `Forward` function (line 12) that takes three parameters. The first is the forwarding table in which the entries are in descending order of prefix length. The second is a `Zen<Header>`, which is an IP header. More generally,

```

1 public class Header {
2     public Ip DstIp;
3     public Ip SrcIp;
4     ....
5 }
6
7 public class Packet {
8     public Header OverlayHeader;
9     public Option<Header> UnderlayHeader;
10 }
11
12 Zen<byte> Forward(FwdTbl t, Zen<Header> h, int i) {
13     if (i >= t.Rules.Length)
14         return 0; // null interface
15     var r = t.Rules[i];
16     return If(Matches(r, h), r.Port, Forward(h, i+1));
17 }
18
19 Zen<bool> Matches(FwdRule r, Zen<Header> h) {
20     var mask = 0xFFFFFFFF << (32 - r.Prefix.Length);
21     return (h.GetDstIp() & mask) == r.Prefix.Address;
22 }

```

Figure 4: Encoding packet forwarding in Zen

the wrapper type `Zen<T>` represents a value of type `T` that is handled by the Zen library and can be either symbolic or concrete. The third parameter is the line number to start matching from.

The function evaluates the header against the forwarding table and returns a `Zen<byte>` representing the output port. It first checks if the line number is beyond the last rule in the forwarding table. If so, no rule applied to the header so it returns 0 (null interface). Otherwise, it gets the current rule and makes a call through Zen to either return the rule's port number if the rule matches the header, otherwise to continue on to rule `i + 1`. Note that the recursive call takes place through C# and not the Zen library.

Matching the header against the rule is similarly implemented as a simple function that checks if the header's destination IP is matched by the forwarding rule prefix. The library overloads operators such as `==` and `&` to work seamlessly over Zen values. We can encode ACLs in a similar manner as forwarding tables since they too are a list of prioritized rules.

A final step is to encode the semantics of tunneling, whose implementation is shown in Figure 5. There are two additional functions to define the effect of encapsulating and decapsulating a packet given a GRE tunnel. The first function `Encap` adds an underlay header using the tunnel's source and destination IP addresses while copying over all other fields. The second function `Decap` simply strips off the top header by replacing it with `Null<Header>()`.

Composing network models. Composing models of network elements with Zen is as simple as writing new functions that call functions defined in earlier models. Suppose we want to model combined (overlay and underlay) treatment of packets at a switch, accounting for forwarding, ACLs, and tunneling. We might write the functions in Figure 6. The two functions take a packet as input along with an interface, and return a value of type `Zen<Option<Packet>>`

```

23 Zen<Packet> Encap(GreTunnel t, Zen<Packet> pkt) {
24   if (t == null) return pkt;
25   var oheader = pkt.OverlayHeader;
26   var uheader = Create<Header>(
27     Create<Ip>(t.SrcIp), Create<Ip>(t.DstIp),
28     oheader.GetDstPort(), oheader.GetSrcPort(),
29     oheader.GetProtocol());
30   return Create<Packet>(oheader, Some(uheader));
31 }
32
33 Zen<Packet> Decap(GreTunnel t, Zen<Packet> pkt) {
34   if (t == null) return pkt;
35   return Create<Packet>(
36     pkt.OverlayHeader, Null<Header>());
37 }

```

Figure 5: Encoding IP GRE tunnels in Zen

```

38 Zen<Option<Packet>> FwdIn(Intf i, Zen<Packet> p) {
39   var allow = Allow(i.AclIn, p);
40   var decap = Decap(i.GreEnd, p);
41   return If(allow, decap, Null<Packet>())
42 }
43
44 Zen<Option<Packet>> FwdOut(Intf i, Zen<Packet> p) {
45   var port = Forward(i.Device, p, 0);
46   var allow = Allow(i.AclOut, p);
47   var encap = Encap(i.GreStart, p);
48   var pktOut = If(allow, encap, Null<Packet>());
49   return If(port == i.Id, pktOut, Null<Packet>());
50 }

```

Figure 6: Modeling the combined (overlay and underlay) treatment of packets being processed at a device.

as output, which is either null if the packet is dropped, or otherwise a new (possibly modified) packet. The first function applies any inbound policy including the ACL and decapsulation, while the second applies outbound policy, including the forwarding table, outbound acl, and any encapsulation.

4 ANALYZING MODELS WITH ZEN

Zen provides a number of ways to analyze network models.

Simulation. Since Zen models are executable—they are simply C# code—simulations performed by tools like Batfish [11] are straightforward. In particular Zen allows users to pass concrete values of type *T* to arguments expecting a value of type *Zen<T>*. For example, to simulate what happens to a given packet entering the network at a given interface, starting with that packet and interface as concrete inputs, we will repeatedly call *FwdIn* and *FwdOut* functions until the packet is dropped or exits the network (along all paths).

Finding (counter) example inputs. Many verification tasks are based on finding an input that leads to an undesirable behavior. Zen enables this primitive using its *Find* method. Suppose we wanted to know if a flow will be delivered along a path. We can write a

```

51 Zen<Option<Packet>> Fwd(Intf[] path, Zen<Packet> p) {
52   Zen<Option<Packet>> x = Some(p);
53   for (int i = 0; i < path.Length; i++) {
54     var intfIn = path[i];
55     var intfOut = path[i + 1];
56     x = If(x.HasValue, FwdIn(intfIn, x.Value), x);
57     x = If(x.HasValue, FwdOut(intfOut, x.Value), x);
58   }
59 }

```

Figure 7: Modeling forwarding along a given path.

```

60 IEnumerable<PathSet>
61 HSA(Intf i, StateSet<Packet> set) {
62   var q = new Queue<PathSet>();
63   q.Enqueue(new PathSet(i, set));
64   while (!q.IsEmpty()) {
65     var path = q.Dequeue();
66     var intfIn = path.Current; // last interface
67     var tin = InboundTransformer(intfIn);
68     var inSet = tin.TransformForward(path.Set);
69     var forwarded = false;
70     foreach (var intfOut in intfIn.Device.Nbrs) {
71       var tout = OutboundTransformer(intfOut);
72       var outSet = tout.TransformForward(inSet);
73       if (outSet.IsEmpty()) continue;
74       forwarded = true;
75       q.Enqueue(path.Extend(intfOut, outSet));
76     }
77     if (!forwarded) yield return path;
78   }
}

```

Figure 8: Implementing HSA using Transformers.

function such as the one in Figure 7 to capture how a flow traverses a path through the network. Zen can then reason about it:

```

79 var f = Function(pkt => Fwd(path, pkt));
80 f.Find((pkt, result) => result.HasValue);

```

The first line creates a *ZenFunction* that the library can manipulate, and the second line asks for a packet that is delivered along the path. Packet delivery is indicated by checking that the result of *f* should have a value. Under the covers, Zen can leverage various forms of symbolic reasoning to find an example (if any) input.

To find if a packet can reach node *A* to *B*, along *any* path, we can iterate over all possible paths between those two nodes. If the execution of *Find* uses SMT-based reasoning, we would have implemented a verifier akin to Anteater [29], though Zen is not limited to just that reasoning method alone.

Computing with sets. Many network analyses reason about sets of objects [4, 5, 22, 41] instead of finding examples. Zen enables such reasoning via *transformers* that can manipulate large sets of objects. For example, we can construct a transformer for the *FwdOut* function for an interface *i*:

```

81 var f = Function(pkt => FwdOut(i, pkt).HasValue);
82 StateSetTransformer<Packet, bool> t = f.Transformer();

```



```

e ::= c | e1 < e2 | e1 + e2 | e1 - e2 | e1 * e2 | e1 & e2 | (e1 | e2) |
    not e | e1 or e2 | e1 and e2 | create[τ](e, ..., e) |
    e.f | e1[f:=e2] | if e1 then e2 else e3 | [] | e1::e2 |
    case e1 of e2 e3 | adapt[τ1, τ2](e)
τ ::= bool | byte | short | ushort | int | uint | long | ulong
    (τ1, τ2) | {f=τ1, ..., f=τn} | List[τ] | Option[τ]

```

Figure 9: Zen abstract language syntax.

Once a user creates a transformer, Zen can automatically compute the `TransformForward` set that represents the set of output objects that correspond to the given input objects and `TransformReverse` set that represents the set of input objects that correspond to the given output objects. These capabilities enable users to build efficient analyses without worrying about the implementation.

Figure 8 shows an implementation of HSA [22], which computes sets of reachable packets from an interface along all paths. It uses the inbound and outbound transformers for network interfaces (built using `FwdIn` and `FwdOut` functions respectively; not shown in the figure) and pushes packet sets through the network to explore all paths. `TransformForward` computes the packet sets at each step.

5 LANGUAGE

The Zen language is designed to be as simple as possible without a priori limiting what users can encode. At its core it is a simple expression language, whose abstract syntax (e) is shown in Figure 9. It supports most logical, arithmetic, and bitwise operations, as well as ways to create objects, get and update their fields (e.f, e₁[f:=e₂]), perform conditional logic (if), and add to (e₁::e₂) and match on lists (case). The types supported by Zen are given by τ. These include primitive types such as an unsigned 32-bit integer (uint) as well as composite types such as tuples, objects, and lists.

To make Zen extensible, we include a special expression type: `adapt[τ1, τ2](e)` that allows for implementing operations over new types by converting them to types that Zen knows how to handle. For instance, Zen currently implements dictionaries by representing them as lists of tuples with the most recent elements at the head of the list, and it implements options by representing them as a class with flag and value fields.

6 IMPLEMENTATION

The Zen framework is implemented in over 15K lines of C# code and available as open source software¹. Zen currently supports several backends. One is for bounded model checking that can use either an SMT (via Z3 [8]) solver or a high-performance binary decision diagram (BDD) solver. For the SMT backend, Zen encodes all primitive operations using the theory of bitvectors before bitblasting [16] the formulas to SAT. Another backend uses the transformer API to perform unbounded model checking and also leverages the BDD backend. Transform operations such as `TransformForward` are implemented using standard pre/post image computation via existential quantification [7]. All the backends in Zen use the reflection capabilities of C# to introspect the types of objects at runtime, and thus build efficient symbolic representations.

¹<https://github.com/microsoft/zen>

Variable ordering heuristics. For the BDD backend, Zen uses a custom analysis, similar to alias analyses in traditional programming languages, to find a strategy for ordering variables. BDDs can often perform very well but are highly sensitive to the order of their variables [33]. For example, when two variables are compared for (in)equality, Zen ensures their orderings will be interleaved, as any other ordering will result in an exponential memory blowup [1]. For instance, in the following function:

```
83 Function<(int,int),int,bool>((x,y) => x.Item2 == y)
```

the second component for x must be interleaved with y.

If Zen detects that different transformers have different variable ordering requirements, it performs a second optimization whereby it allocates a new set of unique variables for the second transformer rather than reusing those for the first. Instead, it converts between the sets of variables dynamically at runtime using a BDD substitution operation. When possible to efficiently order the variables, it avoids this runtime conversion. Moving this translation to runtime in many cases allows for implementing transformers that would otherwise be impossible due to state space explosion.

Composite data structures. To implement complex data structures like lists, Zen uses a variable to represent the list length and another collection of variables to represent the list elements for different sized lengths. It then employs a type-driven merging operation similar to that employed by Rosette [38]. The maximum list length is controlled via an optional parameter to the `Find` function.

7 PRELIMINARY EVALUATION

We present results from preliminary experiments that show the feasibility and promise of our language-based approach. They show that expressing a range of network functionality is easy and the performance overhead of a general solver that is not functionality specific is acceptable.

Expressiveness. To demonstrate Zen’s expressiveness, we write implementations for several networking components such as router ACLs and longest-prefix-based forwarding, as well as control plane route policies offered from commercial vendors and more. For each implementation, we report on the number of lines of code required to model the component using Zen. Table 2 shows the results. In general, we find that the implementations are straightforward and easy to express. Moreover, they compare favorably with existing implementations. For example, Minesweeper [3] implements a similar conversion of route-maps to SMT using 1K lines of code, and Bonsai [4] implements a similar conversion with BDDs using over 1K lines. We do note that our implementation is not 1:1 feature compatible with Minesweeper: we implement certain features such as the full AS path, which Minesweeper does not, and do not implement certain features such as OSPF areas that Minesweeper does. However, the Zen encoding gives both a BDD and SMT backend in the same 75 lines of code.

Performance. We evaluate Zen’s performance on two verification tasks. The first is the time to verify an ACL (a data plane analysis), and the second is the time to verify a route map (a control plane analysis). In both cases, the verifier’s task is to find inputs (data packets or routing messages) that match the last line, which requires analyzing the complete ACL or route map. We generated ACLs and

Network Component	Zen Lines	Existing systems
Access Control Lists	28	>500 [11]
LPM-based Forwarding	18	>900 [22]
Route Map Filters	75	>1000 [3, 4]
IP GRE tunnels	21	

Table 2: Lines of code to express common network functionalities in Zen. The third column shows lines of code for encoding similar functionality in current tools.




Figure 10: Zen microbenchmarks for random ACLs and route maps with line tracking and different solvers.

route maps of different sizes randomly, and we ran both BDD and SMT backends. For the ACL analysis, we also ran Batfish, which performs the same analysis using a hand-optimized, BDD-based encoding. Batfish currently does not support verification of route maps. All experiments were performed on a 8-core Intel i7 machine with 16GB of RAM, and each data point in the graphs is the mean value across 100 runs.

Figure 10 shows the results. For the ACL analysis, we see that Zen’s BDD backend is more efficient than the SMT backend. We also see that this backend performs comparably to the hand-optimized Batfish implementation despite having its encoding generated automatically. Thus, general solvers have the potential to match the performance of custom ones.

For the route map analysis, unlike ACL analysis, we see that the SMT backend performs better than the BDD backend. In general, we have found the SMT backend better for reasoning about data structures such as lists. These results show the value of having access to multiple backends, so users can pick the one that is best for their domain and network. This goal would be almost impossible to achieve with custom encodings as one would have to develop multiple different backends for each functionality.

8 BEYOND MODEL ANALYSIS

While our initial focus with Zen is analyzing network models, it has other important use cases. We briefly discuss two such use cases which we have already prototyped.

Testing implementations. Zen models can become the basis for testing the implementations that they model. Given a Zen function f , $f.GenerateInputs()$ produces test inputs with a high-degree of coverage based on symbolic execution [14]. We can test that these inputs are handled by the implementation as expected. For instance, if we have a model for an ACL, we can generate test packets that match on every single rule in the ACL, and then validate that the

implementation processes each packet as expected. This model-based testing approach has been successfully used before [30, 39]. Zen enables a more modular and systematic way to expand such testing to a broad set of network functionalities.

Synthesizing implementations. Zen models are executable, allowing us to directly generate implementations from them. We can compile any Zen function to a real implementation by simply writing: `f.Compile()`. This instructs Zen to generate C# IL instructions, using the `System.Reflection.Emit` API, which will then be just-in-time compiled to assembly at runtime. The resulting implementation runs efficiently. An implementation extracted in this manner will be in sync with the verified model. This property then becomes the foundation for networks whose implementations are provably correct (modulo compiler bugs).

9 RELATED WORK

Zen builds on two prior threads of research:

Network verification. There has been a long line of research on network verification. These works differ in terms of verification algorithms used as well as the network functionality targeted: stateless dataplanes [19, 21–23, 29, 40, 41], stateful dataplanes (e.g., middleboxes) [32, 43], programmable dataplanes (e.g., Click, P4) [9, 28, 35], distributed routing protocols (e.g., BGP, OSPF) [3–5, 10–12], and centralized control planes [6, 15]. While different domains come with their own challenges, they commonly employ translations to standard verification technologies. Zen aims to abstract away this translation. Even tools that use non-standard or domain-optimized solvers (e.g., HSA [22]) can incorporate such solvers as new backends in Zen, allowing for many models to reap their benefits. Beyond simplifying tool development, Zen also allows for easy composition of network models which is challenging or impossible when different tools are implemented using disparate technologies and APIs.

Intermediate verification languages. Zen draws on prior work on IVLs [24, 27, 34, 37] that aim to simplify verification tasks. It shares many common technologies with these languages. For example, its bounded model checker uses a type-aware merging strategy pioneered by Rosette [37]. However, as shown in Table 1, prior IVLs cannot express many common network analyses. To address this limitation, Zen introduces a new *state set* abstraction that allows for directly manipulating sets of values in user code.

Zen shares the linguistic modeling approach of NV [13]. While NV provides high level abstractions for encoding certain network functionalities (e.g., distributed routing) and analyses, Zen’s abstractions are lower-level and more general. Consequently, it can be used to model a wide range of network functionalities and analyses.

10 CONCLUSION

Verification tools today are implemented as monoliths, mixing together domain semantics, analysis engines and solver technologies. We argue for a compositional approach to network modeling and analysis based on a common intermediate language for expressing domain functionality. This approach can enable rapid construction, composition, and verification of domain-specific models for new network functionality and pave the way for fully-verified networks.

REFERENCES

- [1] A. Aziz, S. Tasiran, and R. K. Brayton. Bdd variable ordering for interacting finite state machines. In *Proceedings of the 31st Annual Design Automation Conference, DAC '94*, page 283–288, New York, NY, USA, 1994. Association for Computing Machinery.
- [2] J. Backes, S. Bayless, B. Cook, C. Dodge, A. Gacek, A. J. Hu, T. Khsai, B. Kocik, E. Kotelnikov, J. Kukovec, S. McLaughlin, J. Reed, N. Rungta, J. Sizemore, M. Stalzer, P. Srinivasan, P. Subotić, C. Varming, and B. Whaley. Reachability analysis for aws-based networks. In I. Dillig and S. Tasiran, editors, *Computer Aided Verification*, pages 231–241, Cham, 2019. Springer International Publishing.
- [3] R. Beckett, A. Gupta, R. Mahajan, and D. Walker. A general approach to network configuration verification. In *Proceedings of the Conference of the ACM Special Interest Group on Data Communication, SIGCOMM '17*, pages 155–168, New York, NY, USA, 2017. ACM.
- [4] R. Beckett, A. Gupta, R. Mahajan, and D. Walker. Control plane compression. In *Proceedings of the 2018 Conference of the ACM Special Interest Group on Data Communication, SIGCOMM '18*, page 476–489, New York, NY, USA, 2018. Association for Computing Machinery.
- [5] R. Beckett, A. Gupta, R. Mahajan, and D. Walker. Abstract interpretation of distributed network control planes. *Proc. ACM Program. Lang.*, 4(POPL), December 2019.
- [6] M. Canini, D. Venzano, P. Perešini, D. Kostić, and J. Rexford. A NICE way to test openflow applications. In *Presented as part of the 9th USENIX Symposium on Networked Systems Design and Implementation (NSDI 12)*, pages 127–140, San Jose, CA, 2012. USENIX.
- [7] E. M. Clarke, T. A. Henzinger, H. Veith, and R. Bloem. *Handbook of Model Checking*. Springer Publishing Company, Incorporated, 1st edition, 2018.
- [8] L. De Moura and N. Bjørner. Z3: An efficient smt solver. In *Proceedings of the Theory and Practice of Software, 14th International Conference on Tools and Algorithms for the Construction and Analysis of Systems, TACAS'08/ETAPS'08*, page 337–340, Berlin, Heidelberg, 2008. Springer-Verlag.
- [9] M. Dobrescu and K. Argyraki. Software dataplane verification. In *11th USENIX Symposium on Networked Systems Design and Implementation (NSDI 14)*, pages 101–114, Seattle, WA, 2014. USENIX Association.
- [10] S. K. Fayaz, T. Sharma, A. Fogel, R. Mahajan, T. Millstein, V. Sekar, and G. Varghese. Efficient network reachability analysis using a succinct control plane representation. In *12th USENIX Symposium on Operating Systems Design and Implementation (OSDI 16)*, pages 217–232, Savannah, GA, November 2016. USENIX Association.
- [11] A. Fogel, S. Fung, L. Pedrosa, M. Walraed-Sullivan, R. Govindan, R. Mahajan, and T. Millstein. A general approach to network configuration analysis. In *12th USENIX Symposium on Networked Systems Design and Implementation (NSDI 15)*, pages 469–483, Oakland, CA, May 2015. USENIX Association.
- [12] A. Gember-Jacobson, R. Viswanathan, A. Akella, and R. Mahajan. Fast control plane analysis using an abstract representation. In *Proceedings of the 2016 ACM SIGCOMM Conference, SIGCOMM '16*, pages 300–313, New York, NY, USA, 2016. ACM.
- [13] N. Giannarakis, D. Loehr, R. Beckett, and D. Walker. Nv: An intermediate language for verification of network control planes. In *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2020*, page 958–973, New York, NY, USA, 2020. Association for Computing Machinery.
- [14] P. Godefroid, N. Klarlund, and K. Sen. Dart: Directed automated random testing. *PLDI '05*, page 213–223, New York, NY, USA, 2005. Association for Computing Machinery.
- [15] A. B. M. Gomes, F. A. M. Alves, R. S. Ferreira, and J. A. M. Nacif. Vericonn: a tool to generate efficient interconnection networks for post-silicon debug. In *2015 16th Latin-American Test Symposium (LATS)*, pages 1–6, March 2015.
- [16] L. Hadarean, K. Bansal, D. Jovanović, C. Barrett, and C. Tinelli. A tale of two solvers: Eager and lazy approaches to bit-vectors. pages 680–695, 07 2014.
- [17] S. Hanks, T. Li, D. Farinacci, and P. Traina. *Generic Routing Encapsulation over IPv4 networks*, 1994 (accessed June, 2020).
- [18] C. Hawblitzel, J. Howell, M. Kapritsos, J. R. Lorch, B. Parno, M. L. Roberts, S. Setty, and B. Zill. Ironfleet: Proving practical distributed systems correct. In *Proceedings of the 25th Symposium on Operating Systems Principles, SOSP '15*, page 1–17, New York, NY, USA, 2015. Association for Computing Machinery.
- [19] A. Horn, A. Kheradmand, and M. Prasad. Delta-net: Real-time network verification using atoms. In *14th USENIX Symposium on Networked Systems Design and Implementation (NSDI 17)*, pages 735–749, Boston, MA, March 2017. USENIX Association.
- [20] K. Jayaraman, N. Bjørner, J. Padhye, A. Agrawal, A. Bhargava, P.-A. C. Bissonnette, S. Foster, A. Helwer, M. Kasten, I. Lee, A. Namdhari, H. Niaz, A. Parkhi, H. Pinnamraju, A. Power, N. M. Raje, and P. Sharma. Validating datacenters at scale. In *Proceedings of the ACM Special Interest Group on Data Communication, SIGCOMM '19*, page 200–213, New York, NY, USA, 2019. Association for Computing Machinery.
- [21] P. Kazemian, M. Chang, H. Zeng, G. Varghese, N. McKeown, and S. Whyte. Real time network policy checking using header space analysis. In *Presented as part of the 10th USENIX Symposium on Networked Systems Design and Implementation (NSDI 13)*, pages 99–111, Lombard, IL, 2013. USENIX.
- [22] P. Kazemian, G. Varghese, and N. McKeown. Header space analysis: Static checking for networks. In *Proceedings of the 9th USENIX Conference on Networked Systems Design and Implementation, NSDI'12*, pages 9–9, Berkeley, CA, USA, 2012. USENIX Association.
- [23] A. Khurshid, X. Zou, W. Zhou, M. Caesar, and P. B. Godfrey. Veriflow: Verifying network-wide invariants in real time. In *Presented as part of the 10th USENIX Symposium on Networked Systems Design and Implementation (NSDI 13)*, pages 15–27, Lombard, IL, 2013. USENIX.
- [24] A. S. Köksal, V. Kuncak, and P. Suter. Constraints as control. In *Proceedings of the 39th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL '12*, page 151–164, New York, NY, USA, 2012. Association for Computing Machinery.
- [25] C. Lattner and V. Adve. Llvm: A compilation framework for lifelong program analysis and transformation. In *Proceedings of the International Symposium on Code Generation and Optimization: Feedback-Directed and Runtime Optimization, CGO '04*, page 75, USA, 2004. IEEE Computer Society.
- [26] K. R. M. Leino. Dafny: An automatic program verifier for functional correctness. In *Proceedings of the 16th International Conference on Logic for Programming, Artificial Intelligence, and Reasoning, LPAR'10*, page 348–370, Berlin, Heidelberg, 2010. Springer-Verlag.
- [27] K. R. M. Leino and P. Rümmer. A polymorphic intermediate verification language: Design and logical encoding. In J. Esparza and R. Majumdar, editors, *Tools and Algorithms for the Construction and Analysis of Systems*, pages 312–327, Berlin, Heidelberg, 2010. Springer Berlin Heidelberg.
- [28] J. Liu, W. Hallahan, C. Schlesinger, M. Sharif, J. Lee, R. Soulé, H. Wang, C. Caşcaval, N. McKeown, and N. Foster. P4v: Practical verification for programmable data planes. In *Proceedings of the 2018 Conference of the ACM Special Interest Group on Data Communication, SIGCOMM '18*, pages 490–503, New York, NY, USA, 2018. ACM.
- [29] H. Mai, A. Khurshid, R. Agarwal, M. Caesar, P. B. Godfrey, and S. T. King. Debugging the data plane with antea. In *Proceedings of the ACM SIGCOMM 2011 Conference, SIGCOMM '11*, pages 290–301, New York, NY, USA, 2011. ACM.
- [30] K. L. McMillan and L. D. Zuck. Formal specification and testing of quic. In *Proceedings of the ACM Special Interest Group on Data Communication, SIGCOMM '19*, page 227–240, New York, NY, USA, 2019. Association for Computing Machinery.
- [31] T. Nelson, C. Barratt, D. J. Dougherty, K. Fisler, and S. Krishnamurthi. The margrave tool for firewall analysis. In *Proceedings of the 24th International Conference on Large Installation System Administration, LISA'10*, page 1–8, USA, 2010. USENIX Association.
- [32] A. Panda, O. Lahav, K. Argyraki, M. Sagiv, and S. Shenker. Verifying reachability in networks with mutable datapaths. In *Proceedings of the 14th USENIX Conference on Networked Systems Design and Implementation, NSDI'17*, pages 699–718, Berkeley, CA, USA, 2017. USENIX Association.
- [33] R. Rudell. Dynamic variable ordering for ordered binary decision diagrams. In *Proceedings of the 1993 IEEE/ACM International Conference on Computer-Aided Design, ICCAD '93*, page 42–47, Washington, DC, USA, 1993. IEEE Computer Society Press.
- [34] S. F. Siegel, M. Zheng, Z. Luo, T. K. Zirkel, A. V. Marianiello, J. G. Edenhofner, M. B. Dwyer, and M. S. Rogers. Civi: the concurrency intermediate verification language. In *SC '15: Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, pages 1–12, 2015.
- [35] R. Stoescu, D. Dumitrescu, M. Popovici, L. Negreanu, and C. Raiciu. Debugging p4 programs with vera. In *Proceedings of the 2018 Conference of the ACM Special Interest Group on Data Communication, SIGCOMM '18*, pages 518–532, New York, NY, USA, 2018. ACM.
- [36] B. Tian, X. Zhang, E. Zhai, H. H. Liu, Q. Ye, C. Wang, X. Wu, Z. Ji, Y. Sang, M. Zhang, D. Yu, C. Tian, H. Zheng, and B. Y. Zhao. Safely and automatically updating in-network acl configurations with intent language. In *Proceedings of the ACM Special Interest Group on Data Communication, SIGCOMM '19*, page 214–226, New York, NY, USA, 2019. Association for Computing Machinery.
- [37] E. Torlak and R. Bodik. Growing solver-aided languages with rosette. In *Proceedings of the 2013 ACM International Symposium on New Ideas, New Paradigms, and Reflections on Programming & Software, Onward! 2013*, page 135–152, New York, NY, USA, 2013. Association for Computing Machinery.
- [38] E. Torlak and R. Bodik. A lightweight symbolic virtual machine for solver-aided host languages. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI '14*, page 530–541, New York, NY, USA, 2014. Association for Computing Machinery.
- [39] M. Utting and B. Legeard. *Practical Model-Based Testing: A Tools Approach*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2006.
- [40] G. G. Xie, Jibin Zhan, D. A. Maltz, Hui Zhang, A. Greenberg, G. Hjaltmysson, and J. Rexford. On static reachability analysis of ip networks. In *Proceedings IEEE 24th Annual Joint Conference of the IEEE Computer and Communications Societies*, volume 3, pages 2170–2183 vol. 3, March 2005.
- [41] H. Yang and S. S. Lam. Real-time verification of network properties using atomic predicates. *IEEE/ACM Trans. Netw.*, 24(2):887–900, April 2016.
- [42] H. Yang and S. S. Lam. Scalable verification of networks with packet transformers using atomic predicates. *IEEE/ACM Transactions on Networking*, 25(5):2900–2915,

- 2017.
- [43] Y. Yuan, S.-J. Moon, S. Uppal, L. Jia, and V. Sekar. Netsmc: A custom symbolic model checker for stateful network verification. In *17th USENIX Symposium on Networked Systems Design and Implementation (NSDI 20)*, pages 181–200, Santa Clara, CA, February 2020. USENIX Association.