


Survey in Network Config Analytics

siiba

General background

Debugging a network is hard !

- Complexity of the multiple protocol interactions
- Low level configuration
- Failures

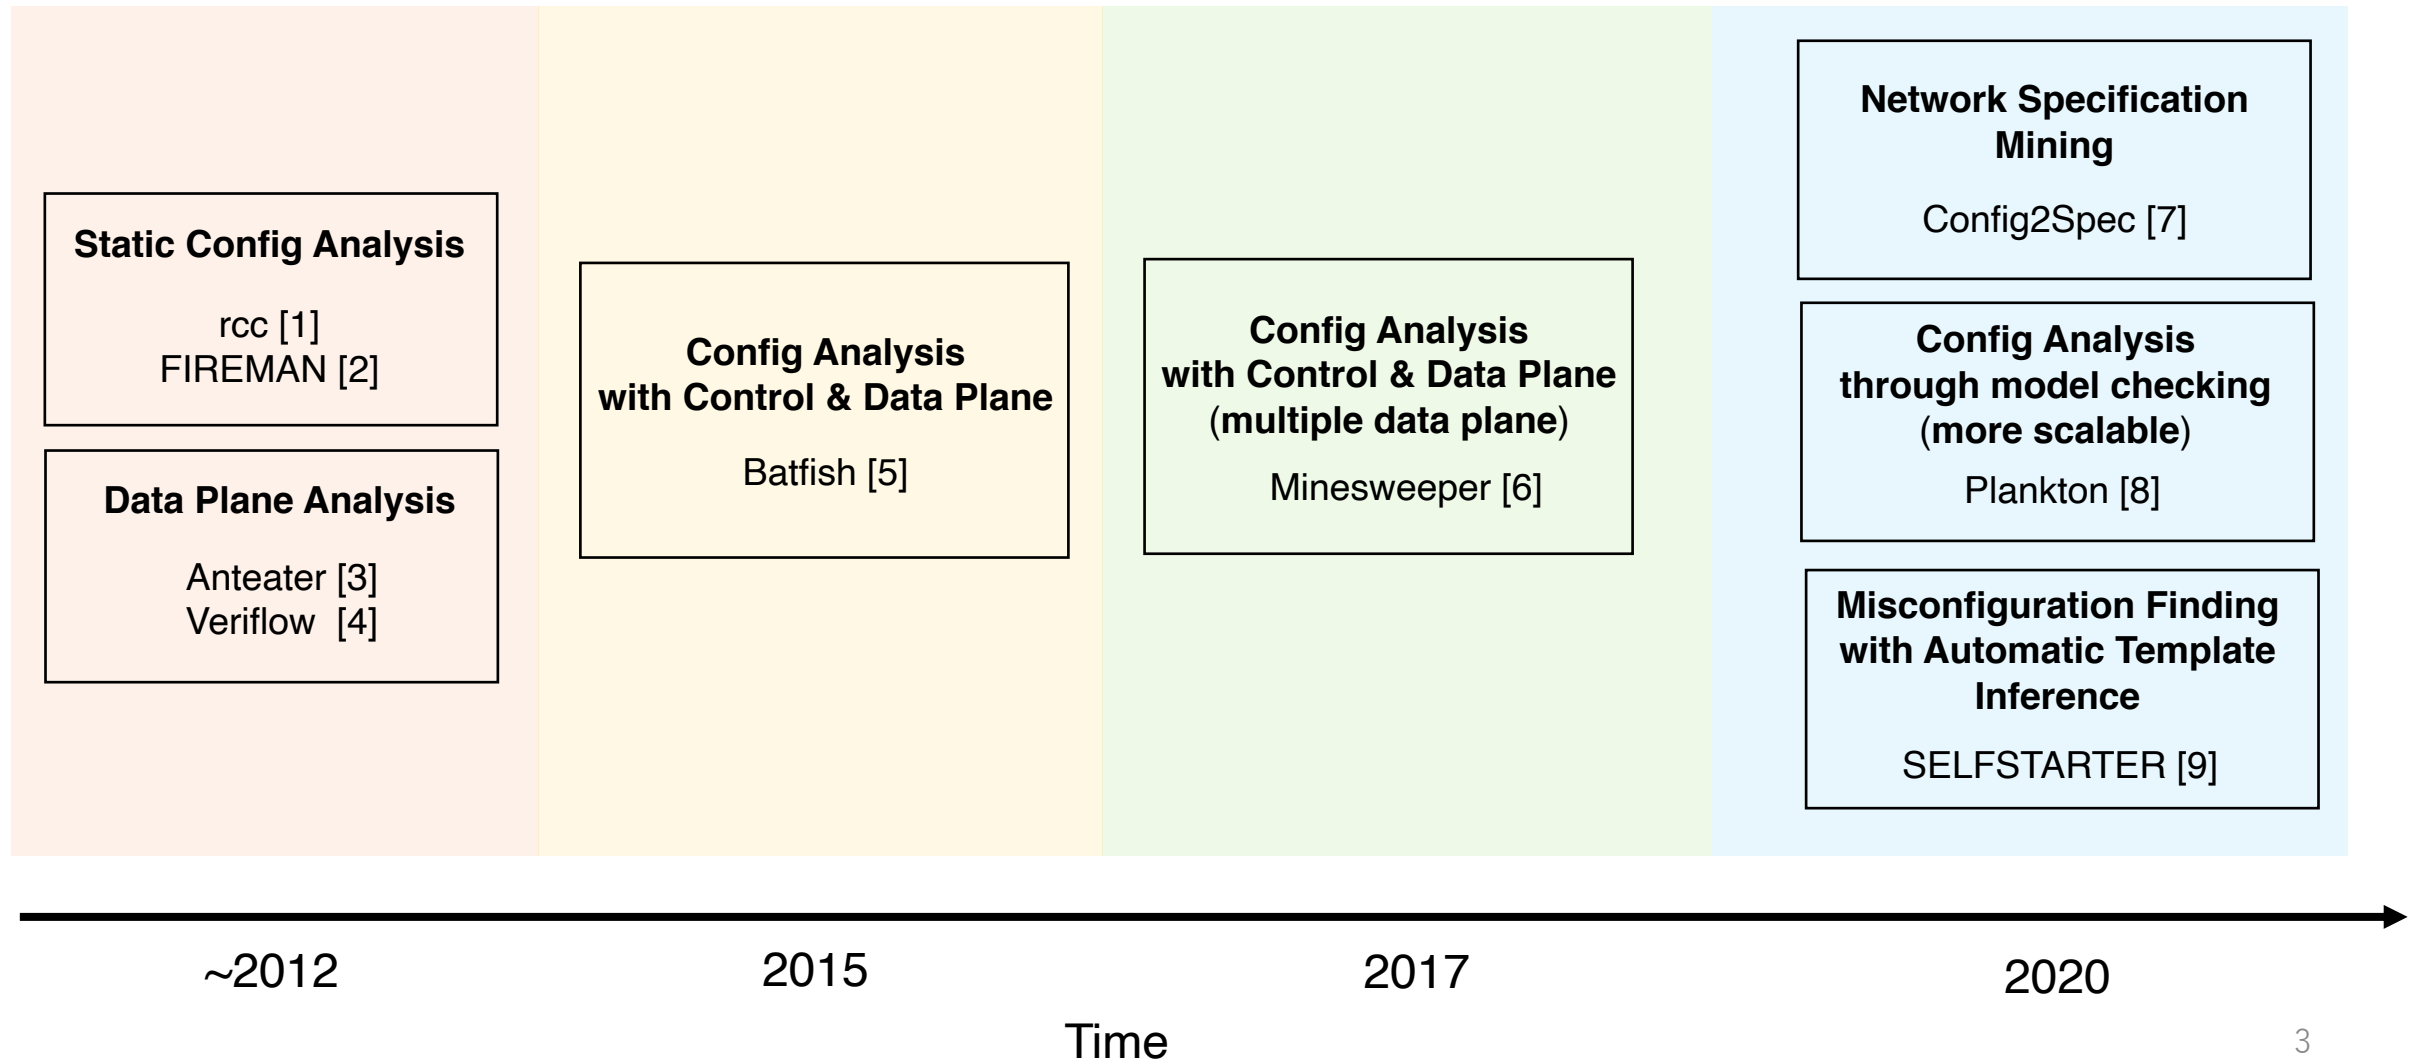


Misconfiguration of the network is so expensive !

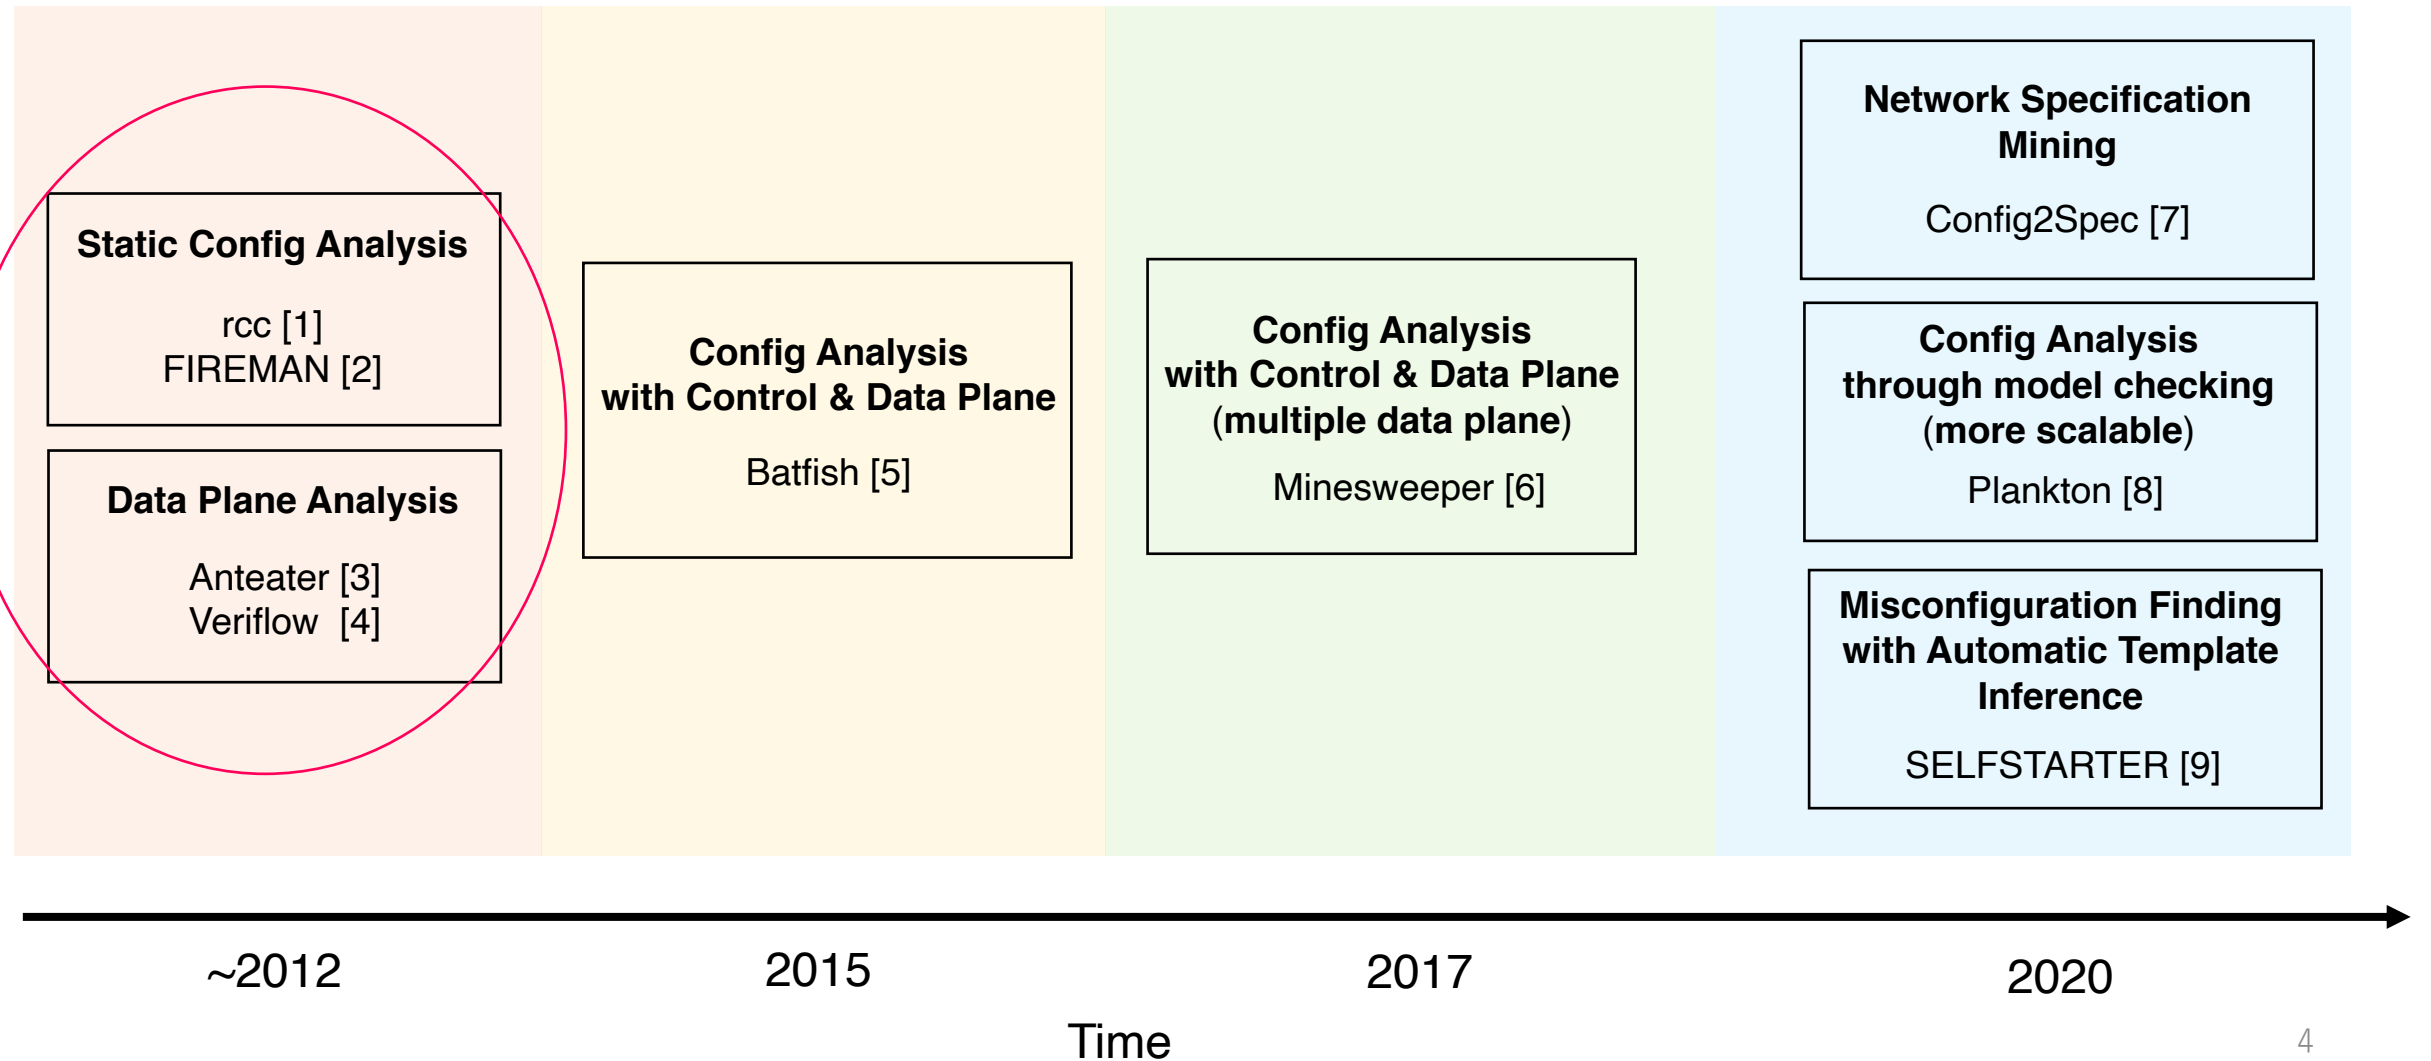
- cost hundreds of thousands of dollars for every hour of downtime



Historical progress



Historical progress



Static Configuration Analysis (~2012s)

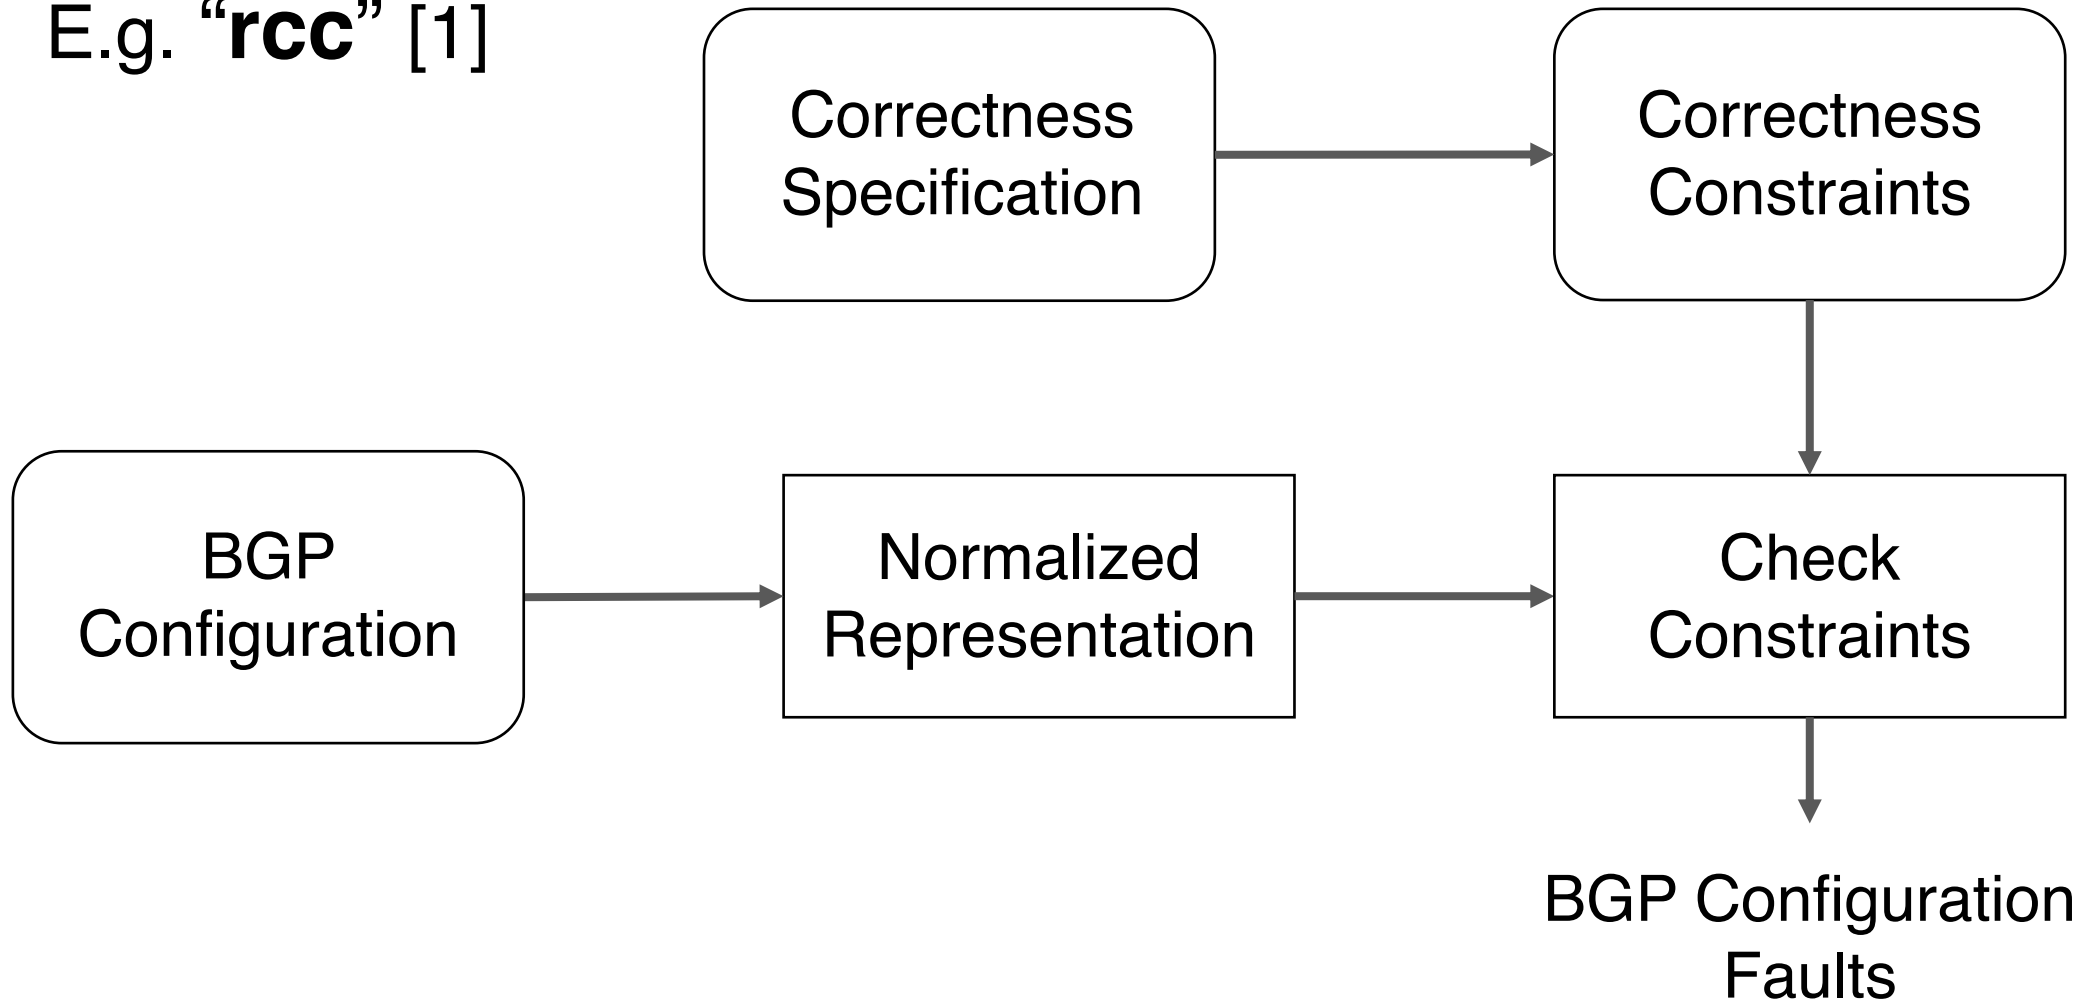
Directly analyzing network configuration files

- + Detecting configuration errors proactively !
- + Can do "what-if" !



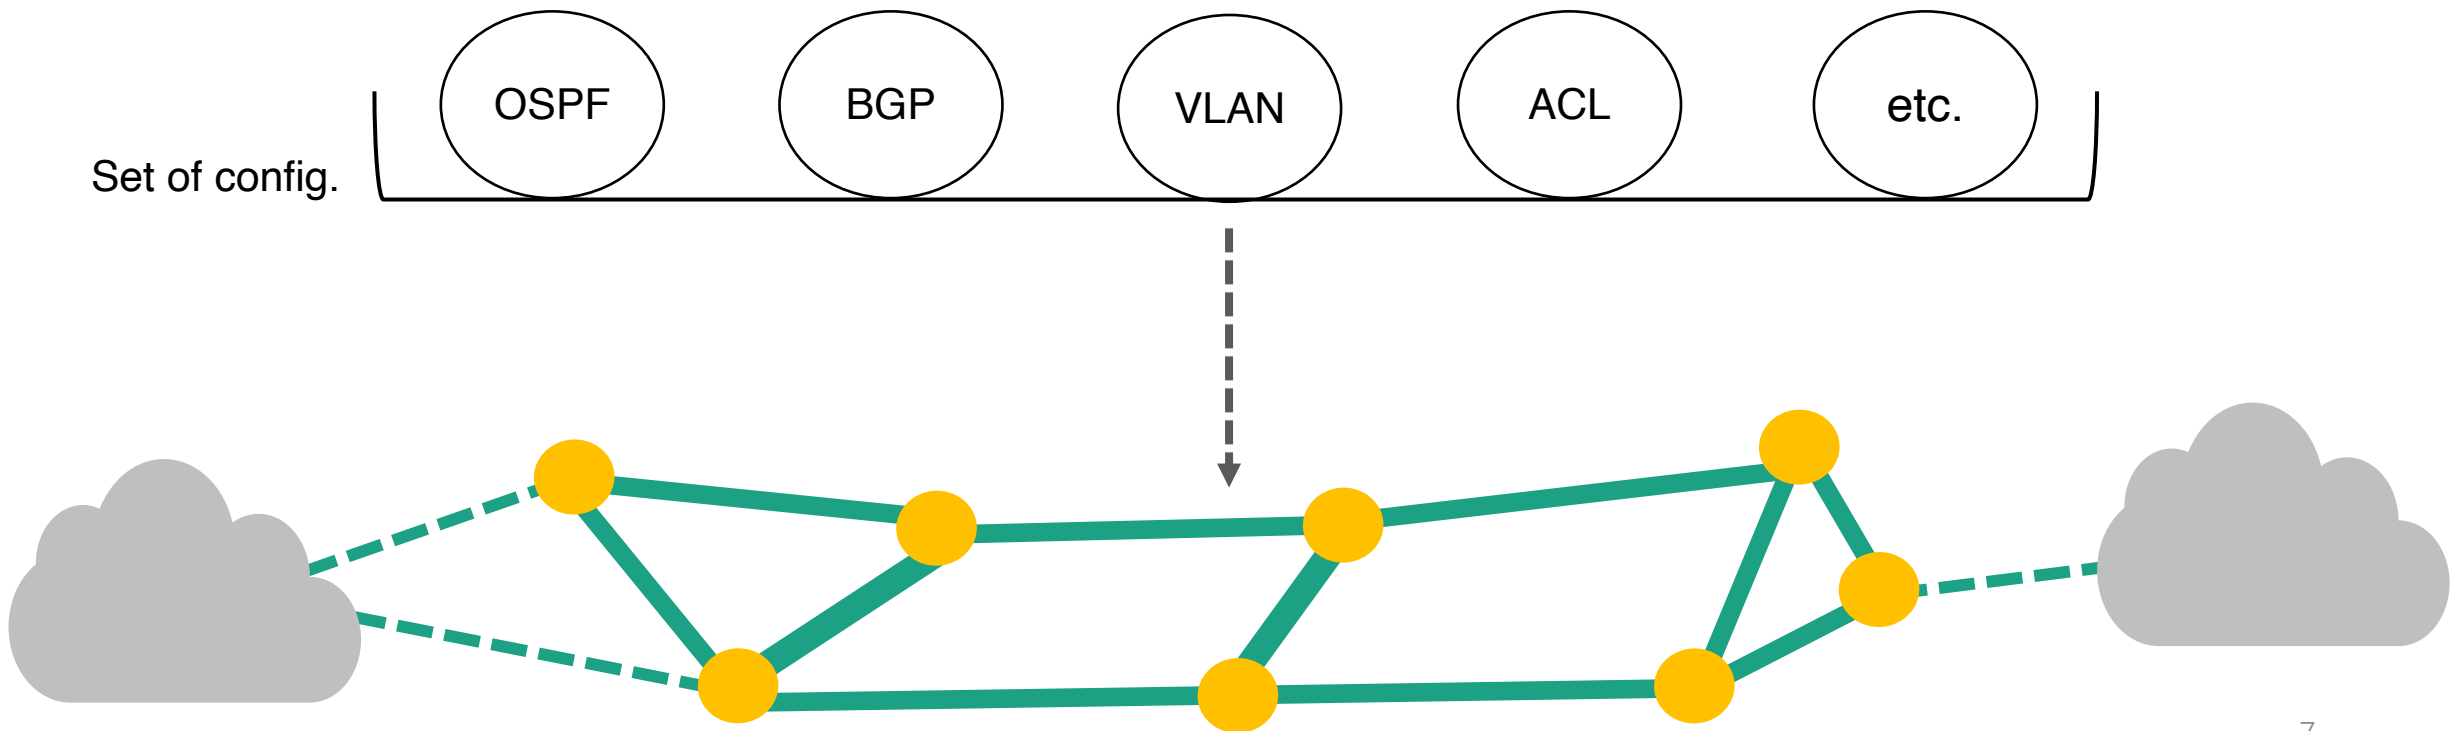
Static Configuration Analysis (~2012)

E.g. “**rcc**” [1]



Static Configuration Analysis (~2012)


- Customized model for specific aspects of config. or properties
 - Unable to handle interactions of the many protocols



Data Plane Analysis (~2012)

Directly analyzing network data plane (Anteater[3], Veriflow[4])

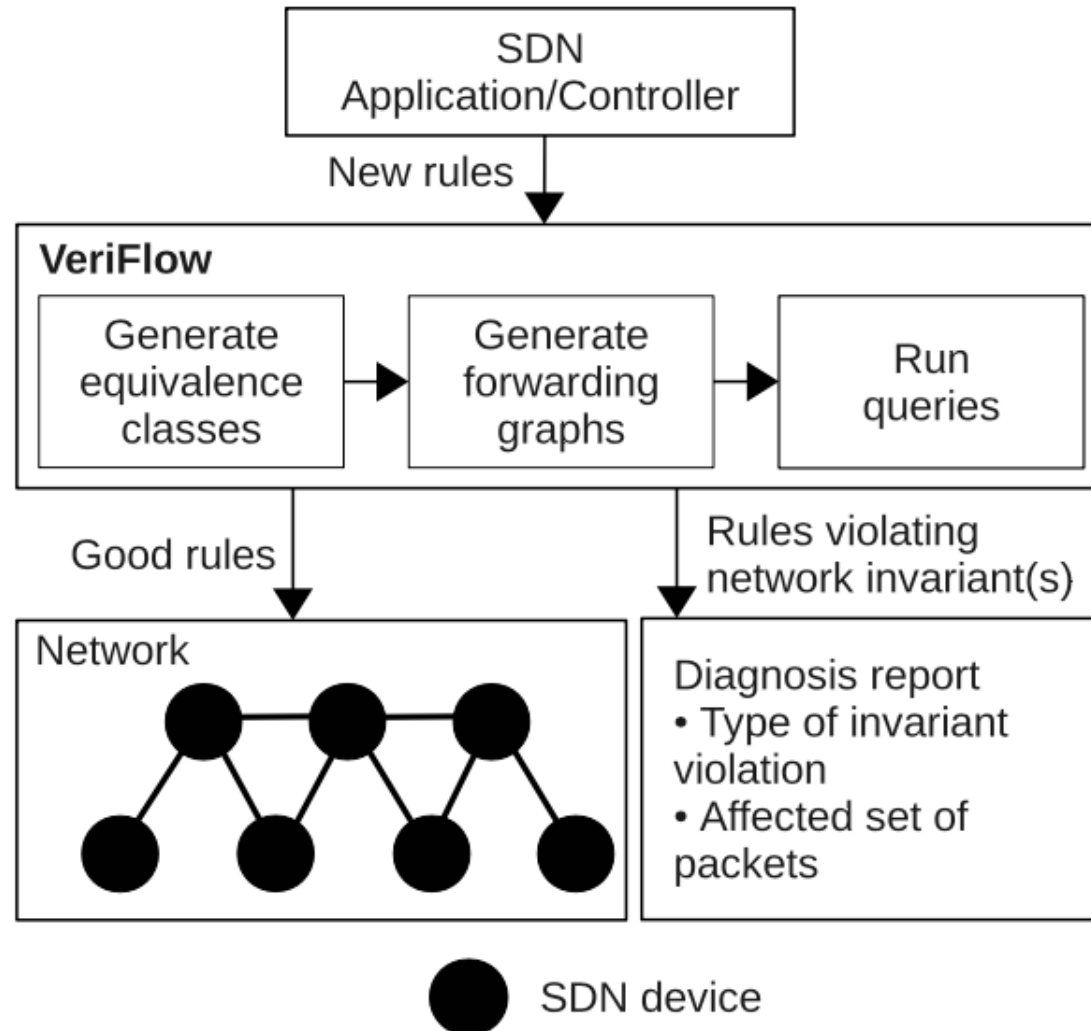
- + Detecting configuration error precisely !
- + Good structure for encoding in various logic !
 - scalable checking with constraint solvers



Data Plane Analysis (~2012)


E.g. “**Veriflow**” [4]

A proxy between
a SDN controller and device



Data Plane Analysis (~2012s)

- Can not prevent the errors proactively
- Need to localize the responsible snippets of configuration

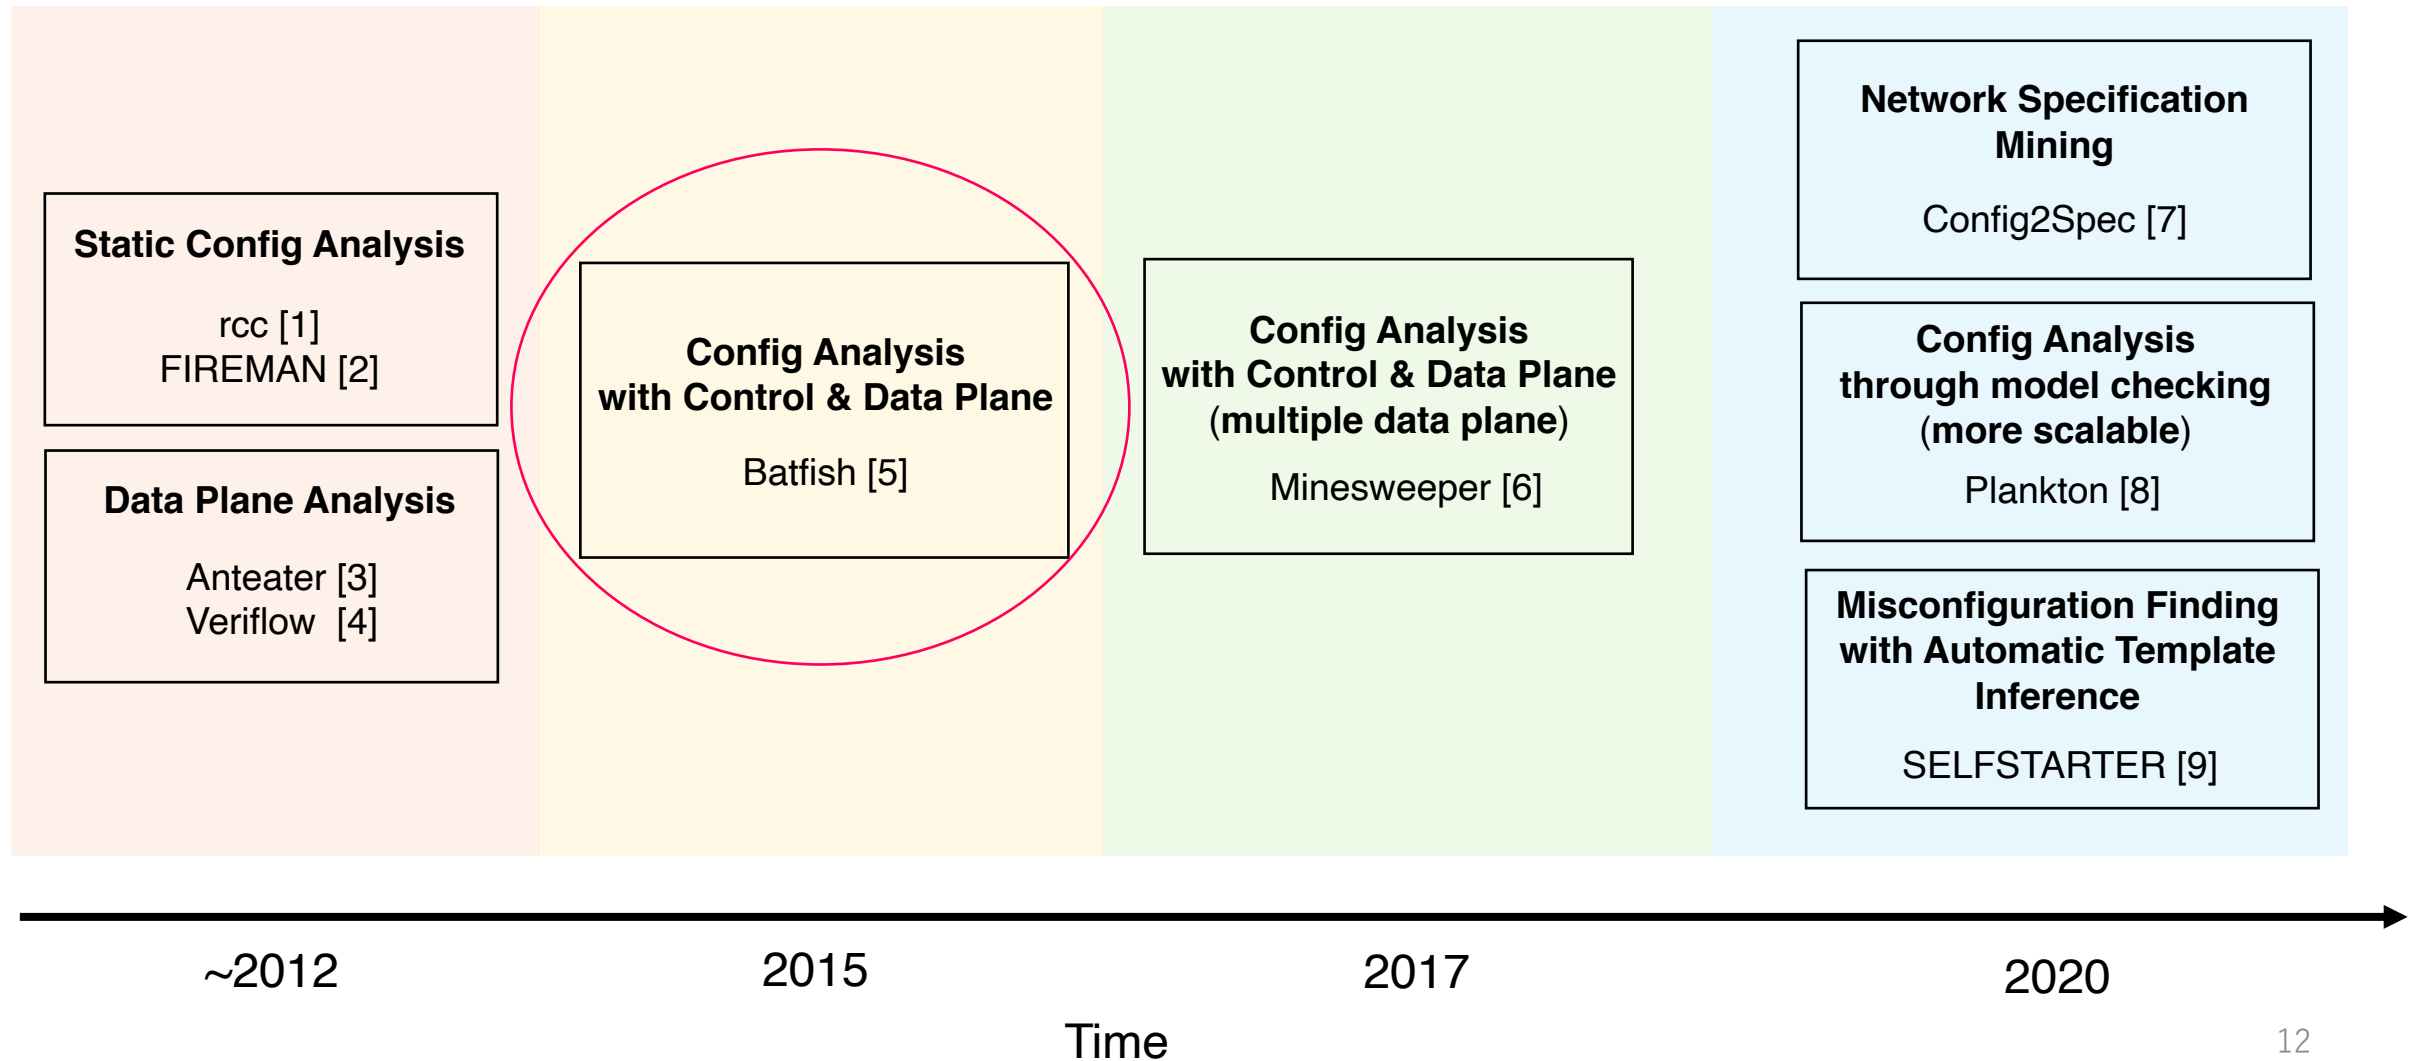


Summary (~2012)

- The approaches have both advantage and disadvantage


	Proactive error analysis	Interaction of multiple protocols
Static configuration analysis	○	×
Data plane analysis	×	○

Historical progress



Background and Motivation


- Can we achieve proactive error analysis and checking of any forwarding property simultaneously?



Approach

Combining the strength of the two prior approaches

1. Deriving a data plane model given a config. and environment
2. Checking the correctness properties with data-plane analysis



Approach

- Combining the strength of the two approaches

	Proactive error finding	Interaction of multiple protocols
Static configuration analysis (~2013)	○	✗
Data plane analysis (~2013)	✗	○
Batfish	○	○

Approach

Stage 1: Configuration file + Topology → Control Plane Model

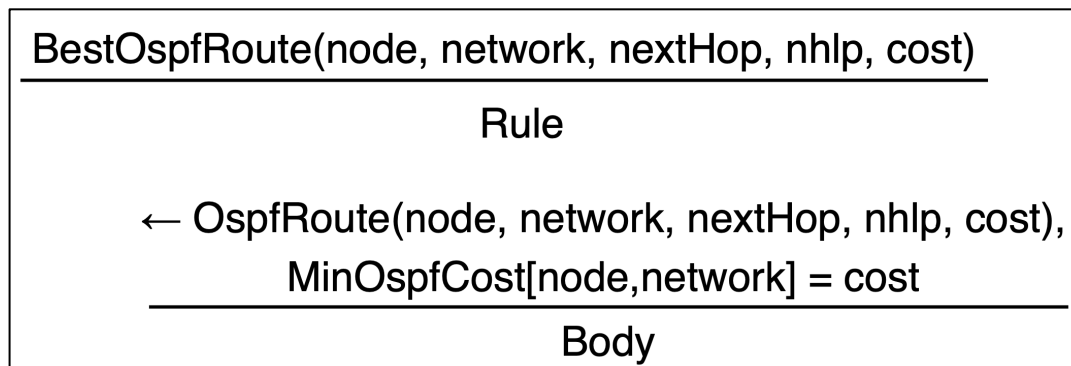
Stage 2: CP. Model + Environment → Data Plane Model

Stage 3: DP. Model + Safety Property → Counterexample


Stage 4: DP. Model + User Input + Counterexample
→ Misconfiguration line

Stage 1: Generating control plane model

- Transforming the configuration with the network topology into a control plane model
- The model is defined in LogiQL



One of the example: A best route of OSPF

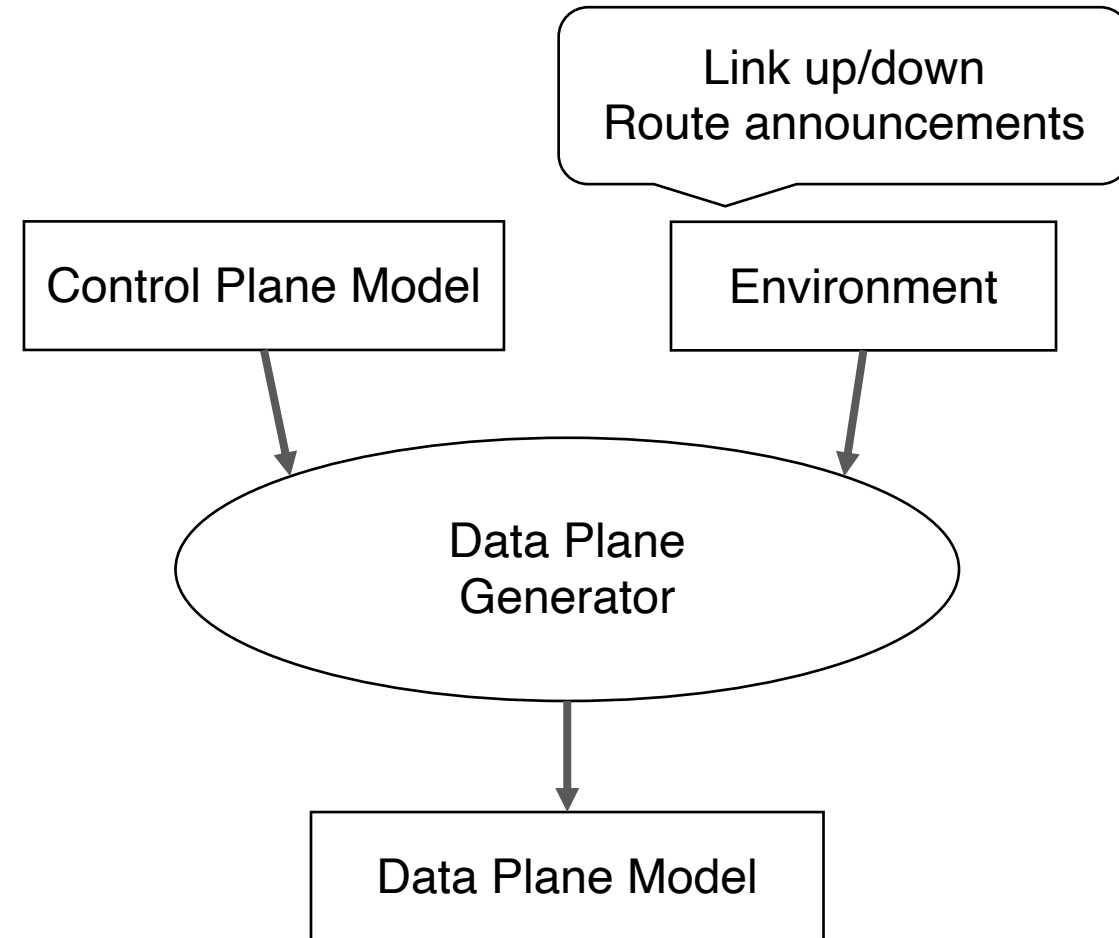


Stage 2: Generating data plane model

- Transforming the control plane model and environment into a data plane model
- The model includes the forwarding behavior as logical facts


Drop(node, flow)
Forward(node, flow, neighbor)

One of the example: Drop and Forward



Stage 3: Checking safety properties

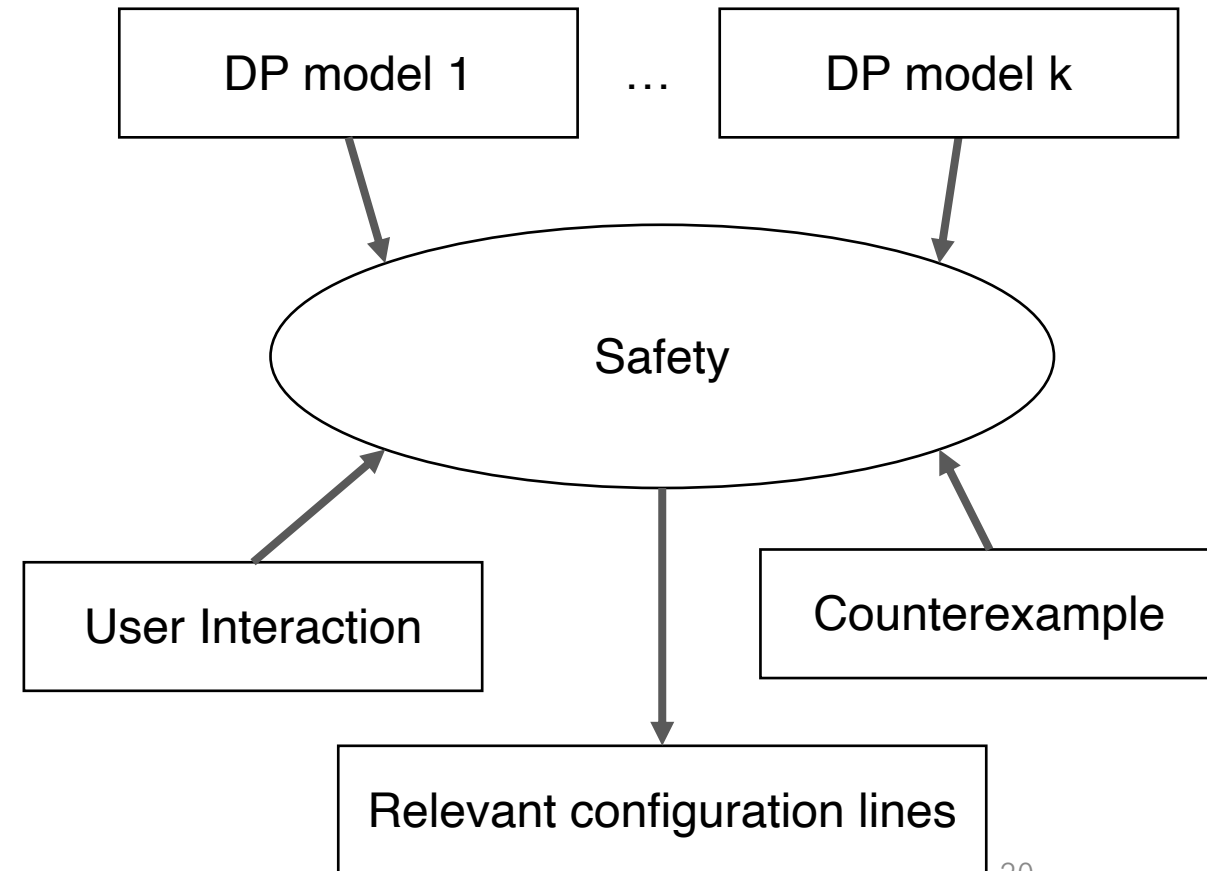
- Translating the data-plane relations and the correctness property to the language of the Z3 constraint solver



Stage 4: Finding error-relevant configuration lines

- Searching the facts in the following order

Counterexample → DP model
→ CP model. → Configuration lines



Evaluation

Net1: 21 routers, 52 AS, Net2: 17 routers, 1AS

Computation time

- Net 1: 238 min, Net 2: 37min

Error detection



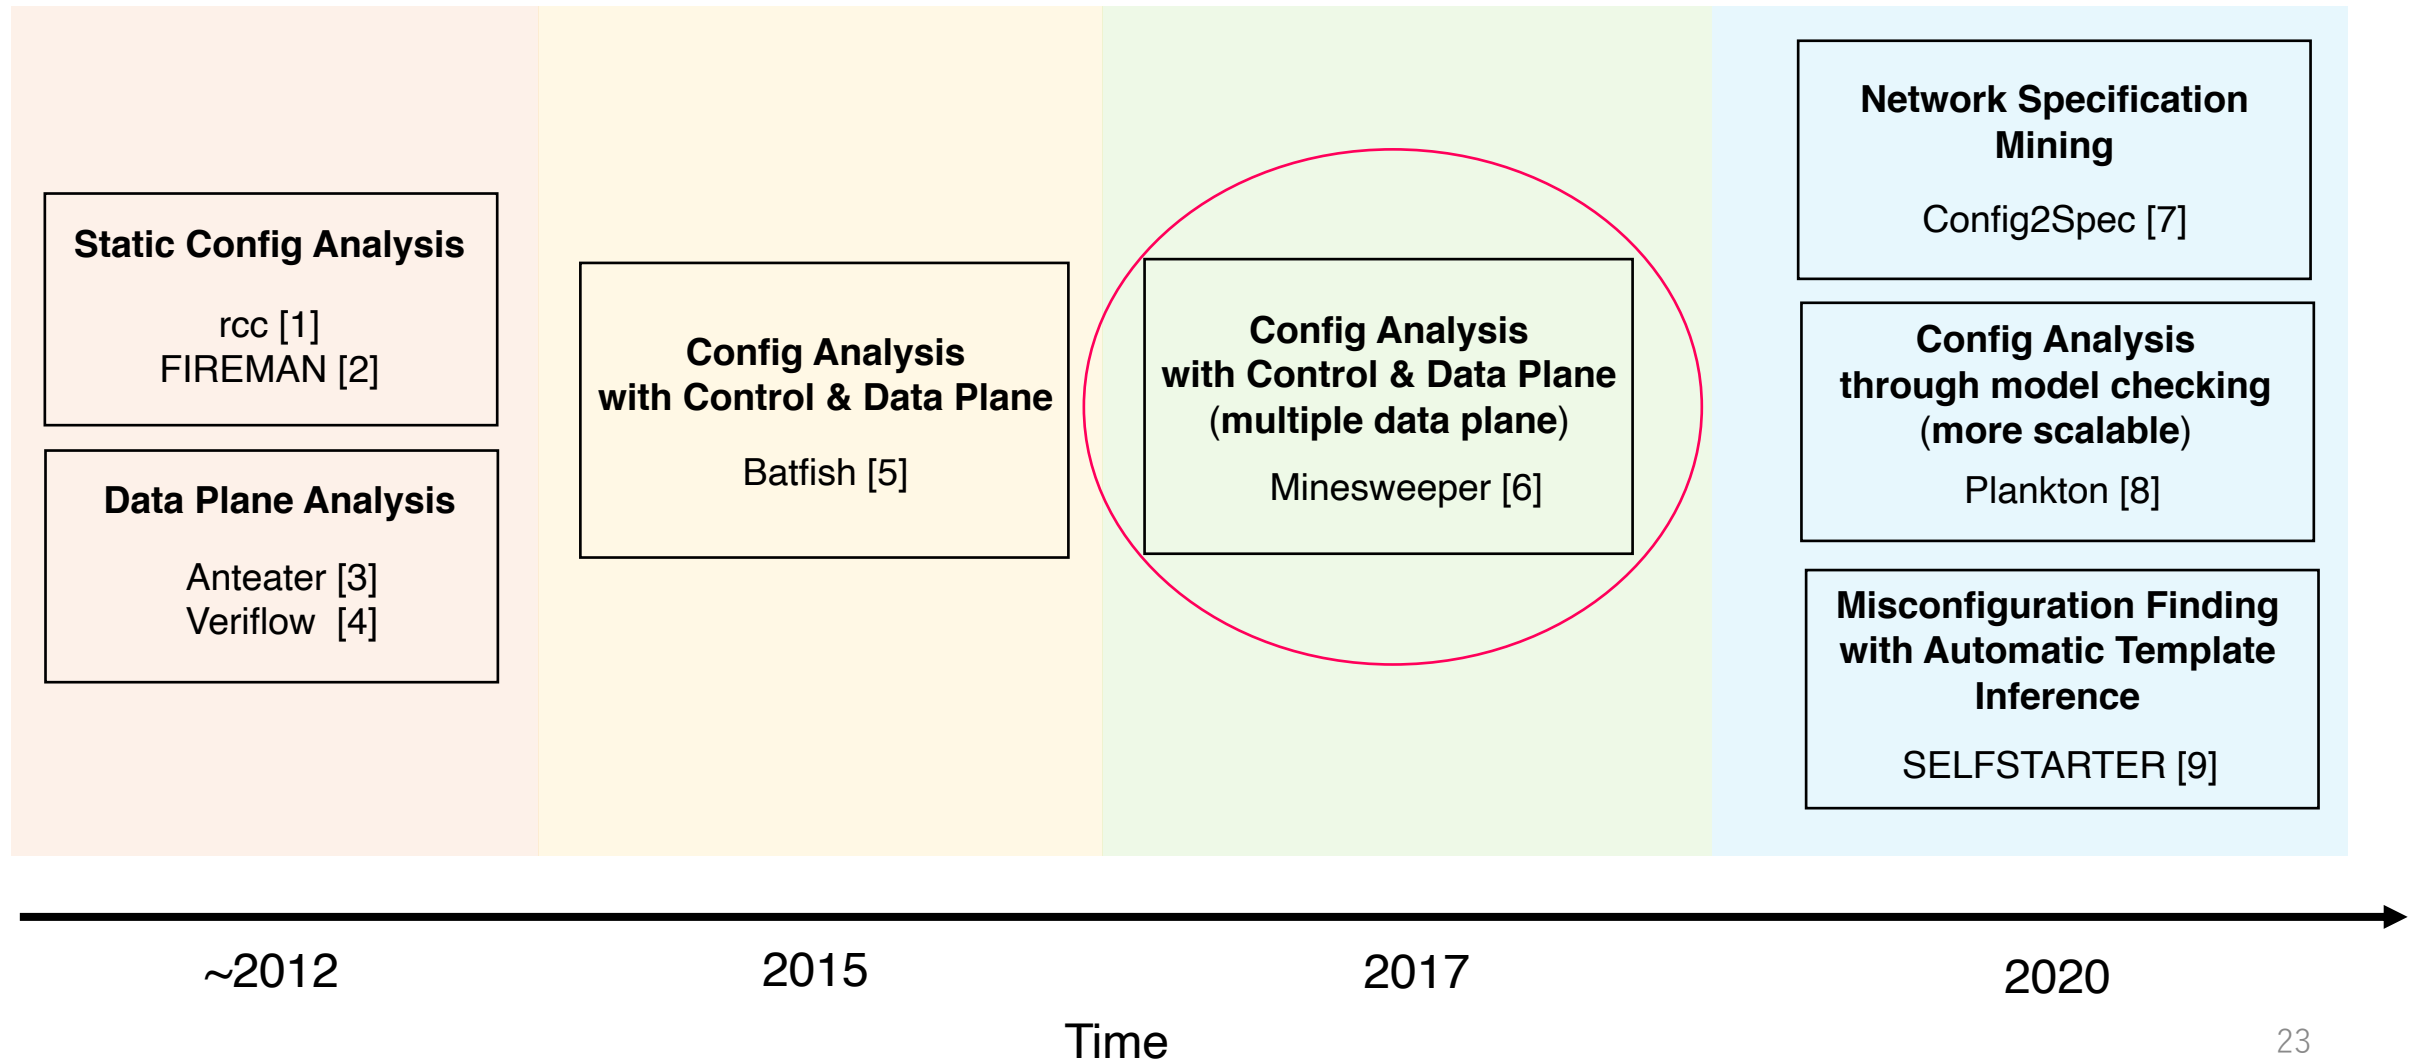
		Total violations	Undesired behaviors	Fixed violations
Net1	Multipath	32 (4)	32 (4)	21 (3)
	Failure	16 (7)	3 (2)	0 (0)
	Destination	55 (6)	55 (6)	1 (1)
Net2	Multipath	11 (3)	11 (3)	11 (3)
	Failure	77(26)	18(7)	0(0)

Summary


- Batfish combines the benefits of prior works to achieve both
proactive analysis & any forwarding property checking !

	Proactive error finding	Interaction of multiple protocols
Static configuration analysis (~2013)	○	✗
Data plane analysis (~2013)	✗	○
Batfish	○	○

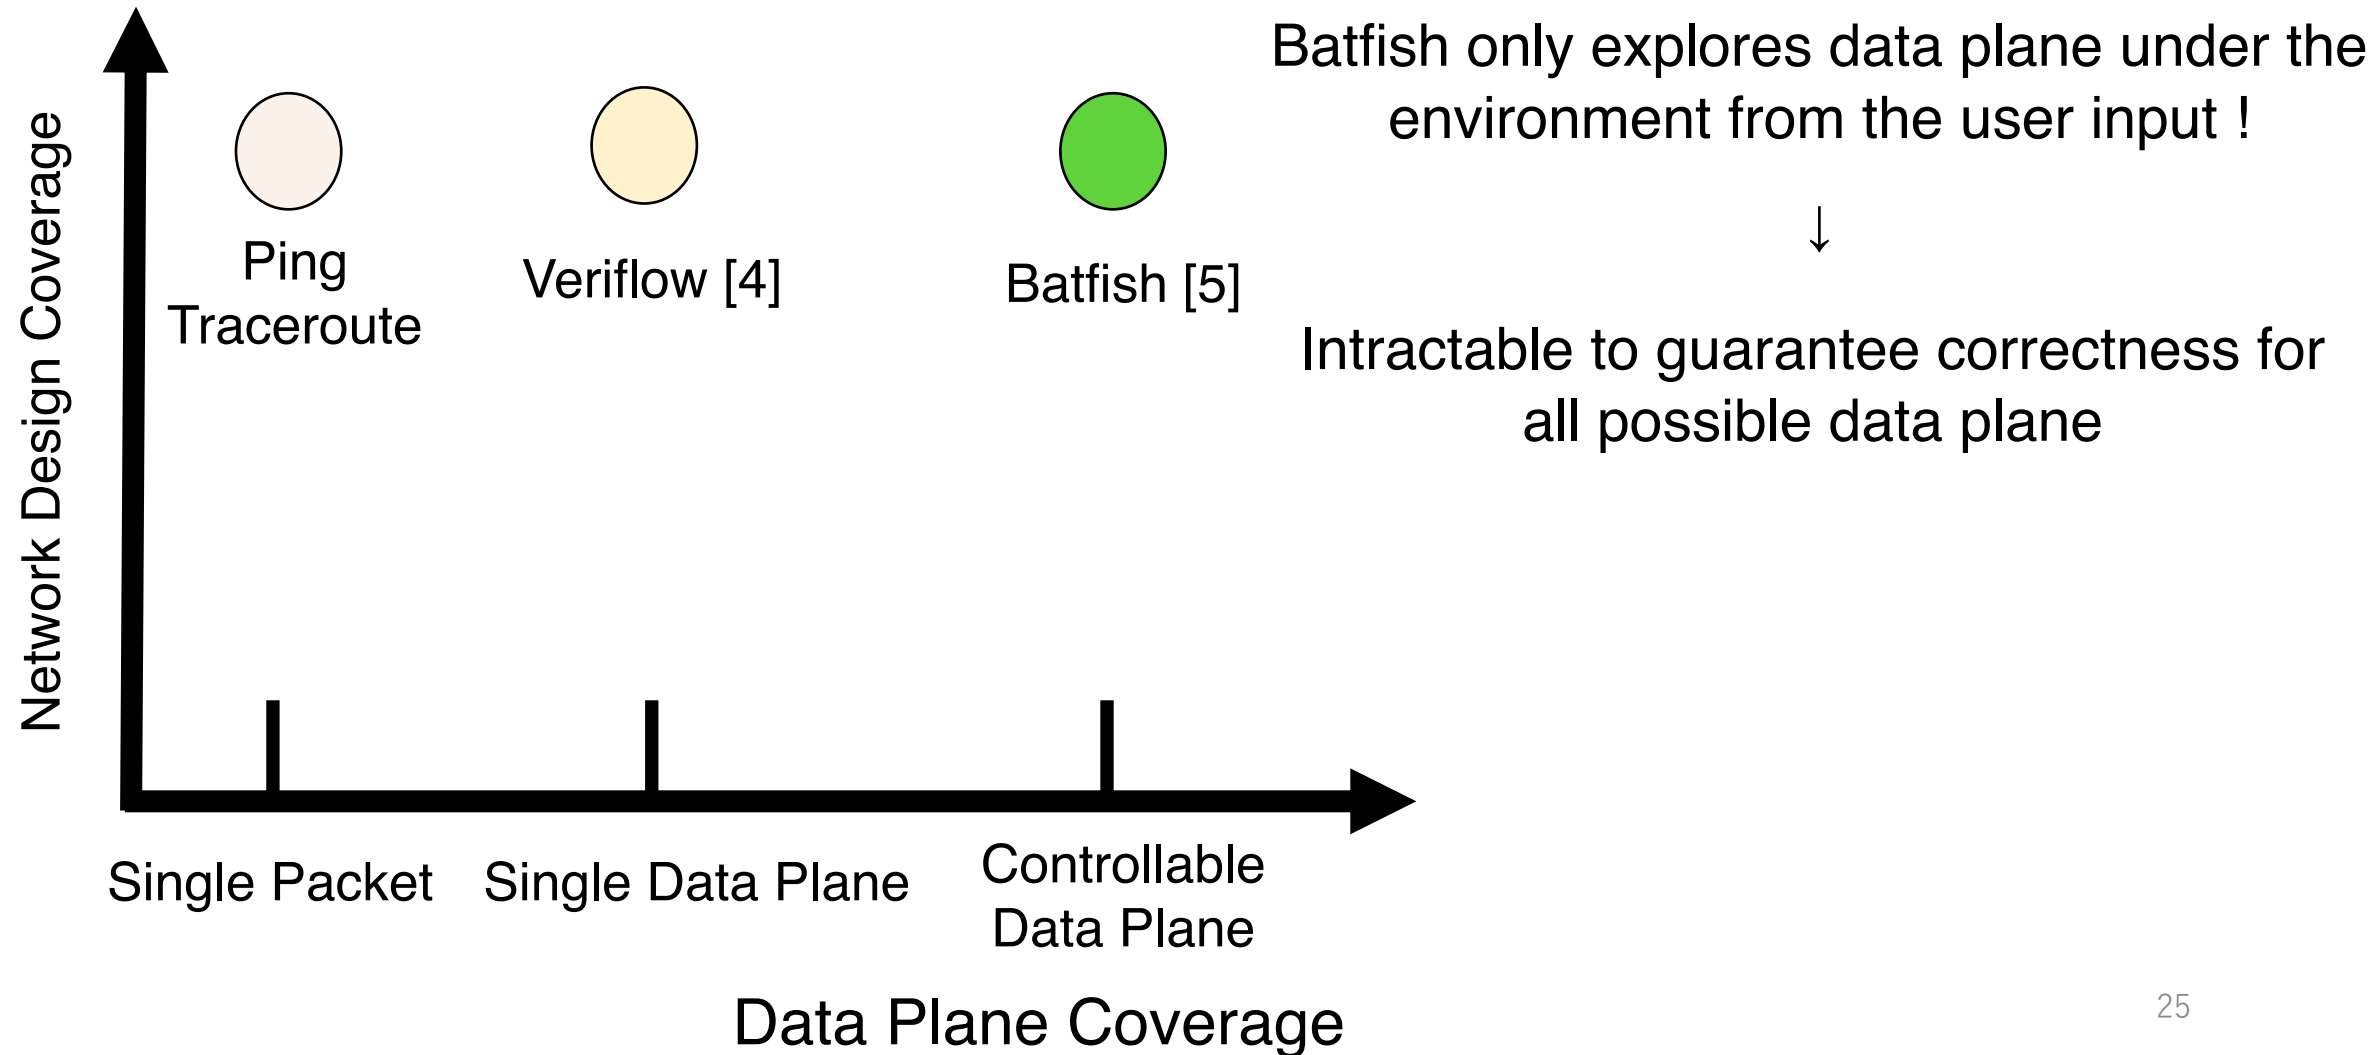
Historical progress



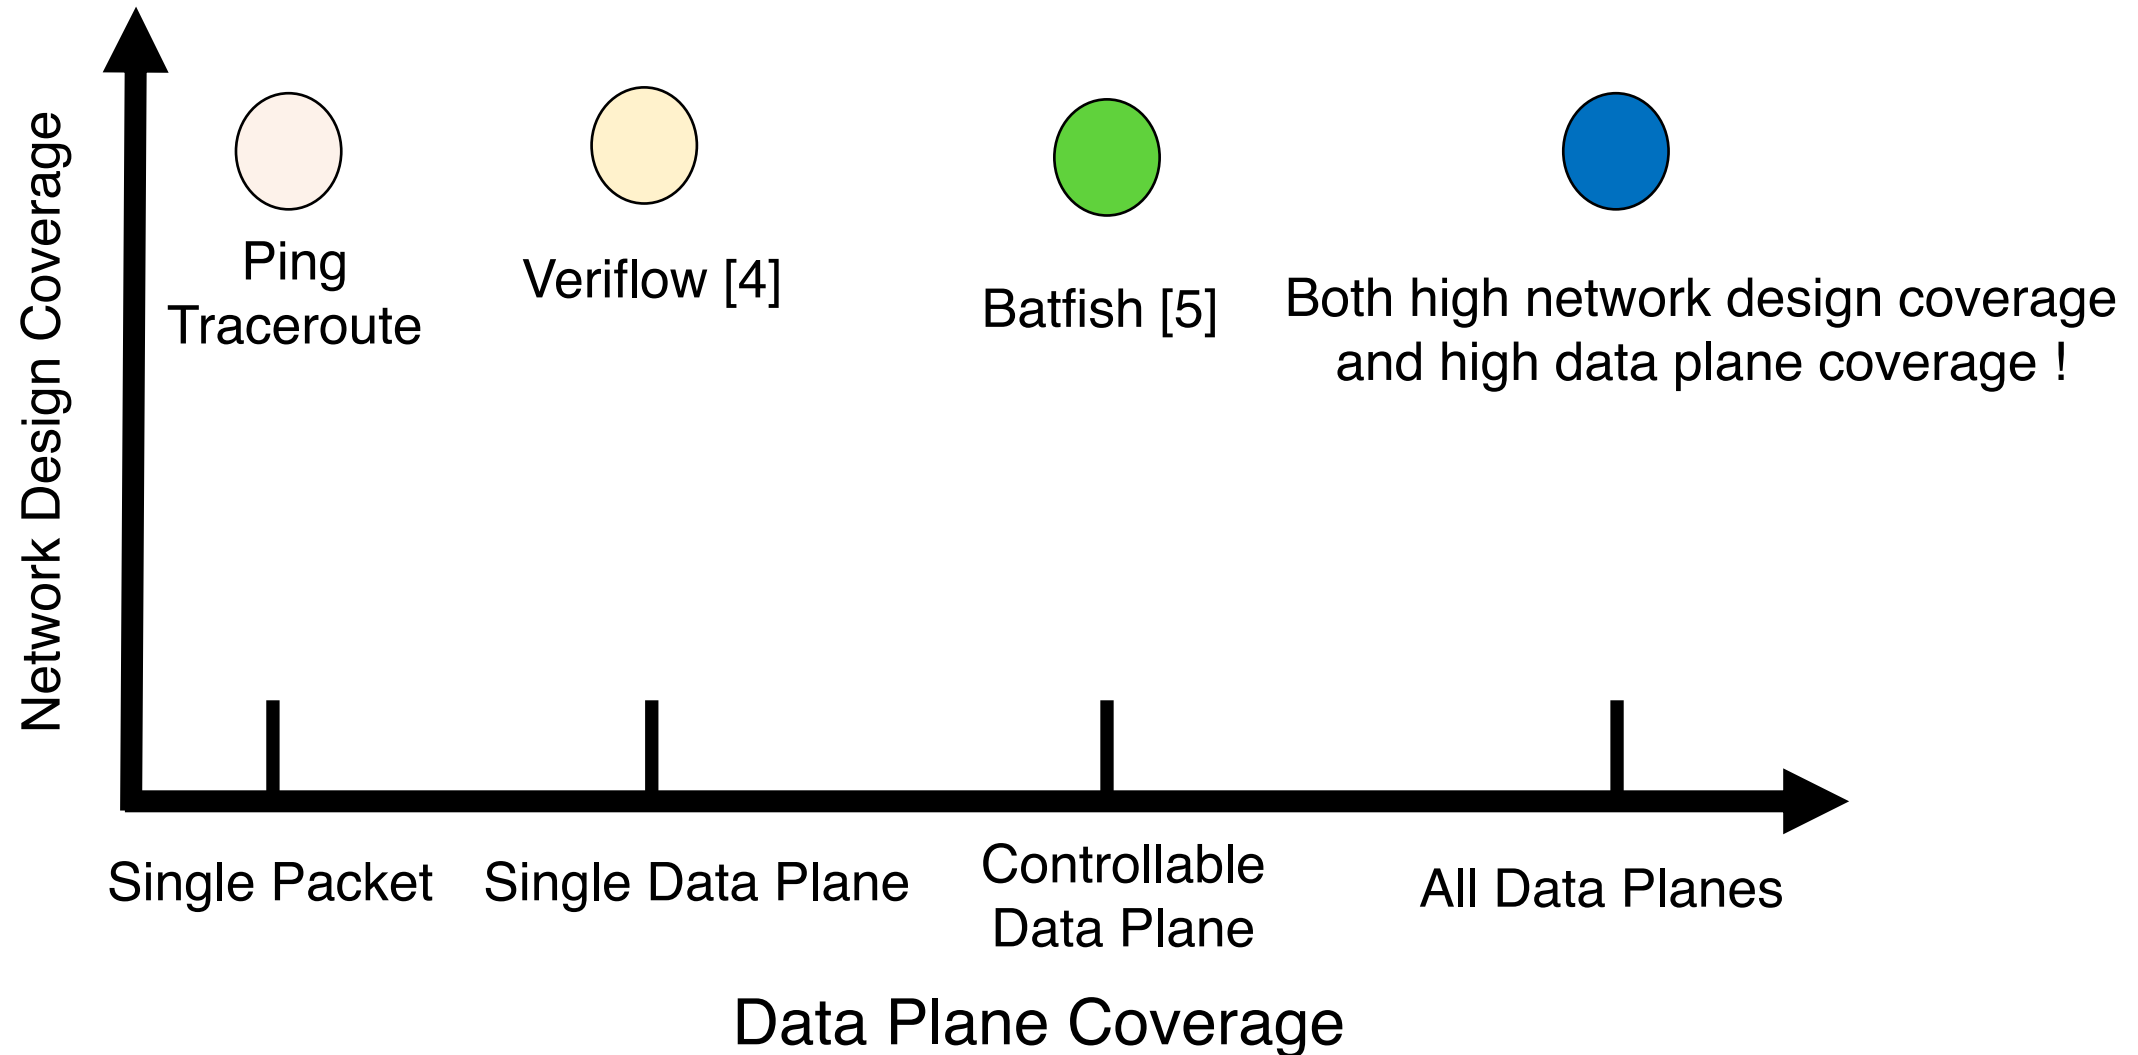
Progress in Network verification




Progress in Network verification



Motivation



Motivation




Minesweeper's key idea

- Encoding the network as a collection of logical constraints
 - Solving the constraints leveraging the off-the-shelf solvers
- Can check many properties for all data planes !

Minesweeper workflow

1. Vendor-Specific Configs



Parse



Encode




3. Constraint Encoding

$192.0.0.0 \leq \text{out.prefix}$
 $\text{out.prefix} \leq 192.1.0.0$
 $\text{best.valid} \rightarrow \text{out.lp} = 120$

+ Property

2. Vendor-Independent Format

Set of configuration



4. Output

Solve




```
batfish -- java -allnone -runmode interactive -- 84x25
Counterexample Found (as2border1<-->as2border2):
Packet:
-----
dstIp: 1.0.0.0
srcIp: 2.0.0.0

Environment Messages:
-----
as2border1, FastEthernet0/0 (BGP):
  community as1_community:
    prefix: 0.0.0.0/1
    protocol metric: 1
as2border2, FastEthernet0/0 (BGP):
  community as1_community:
    prefix: 0.0.0.0/1
    protocol metric: 1

Final Forwarding:
-----
as2border1, FastEthernet0/0 --> ...


batfish>
```

Performance: Scalability



Management interface reachability

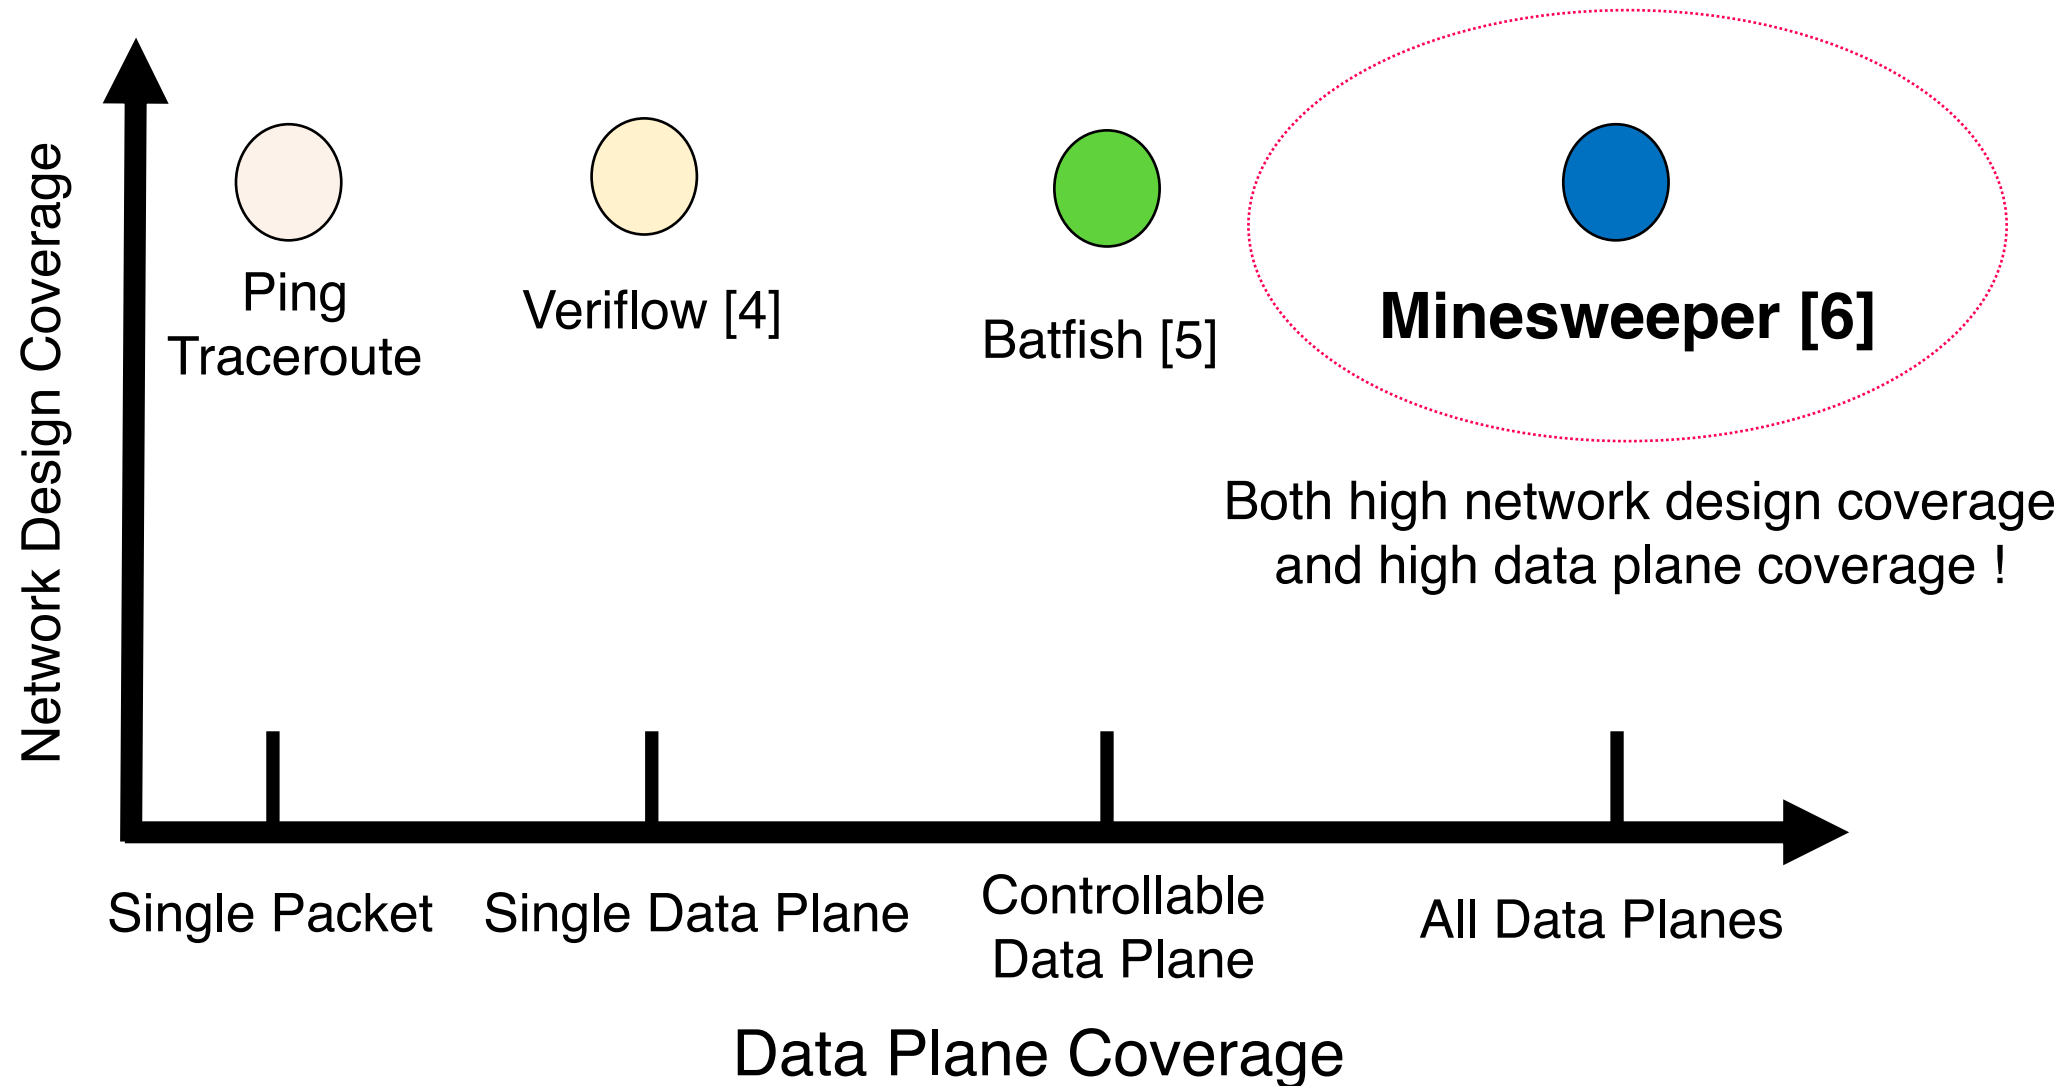
less than 60ms



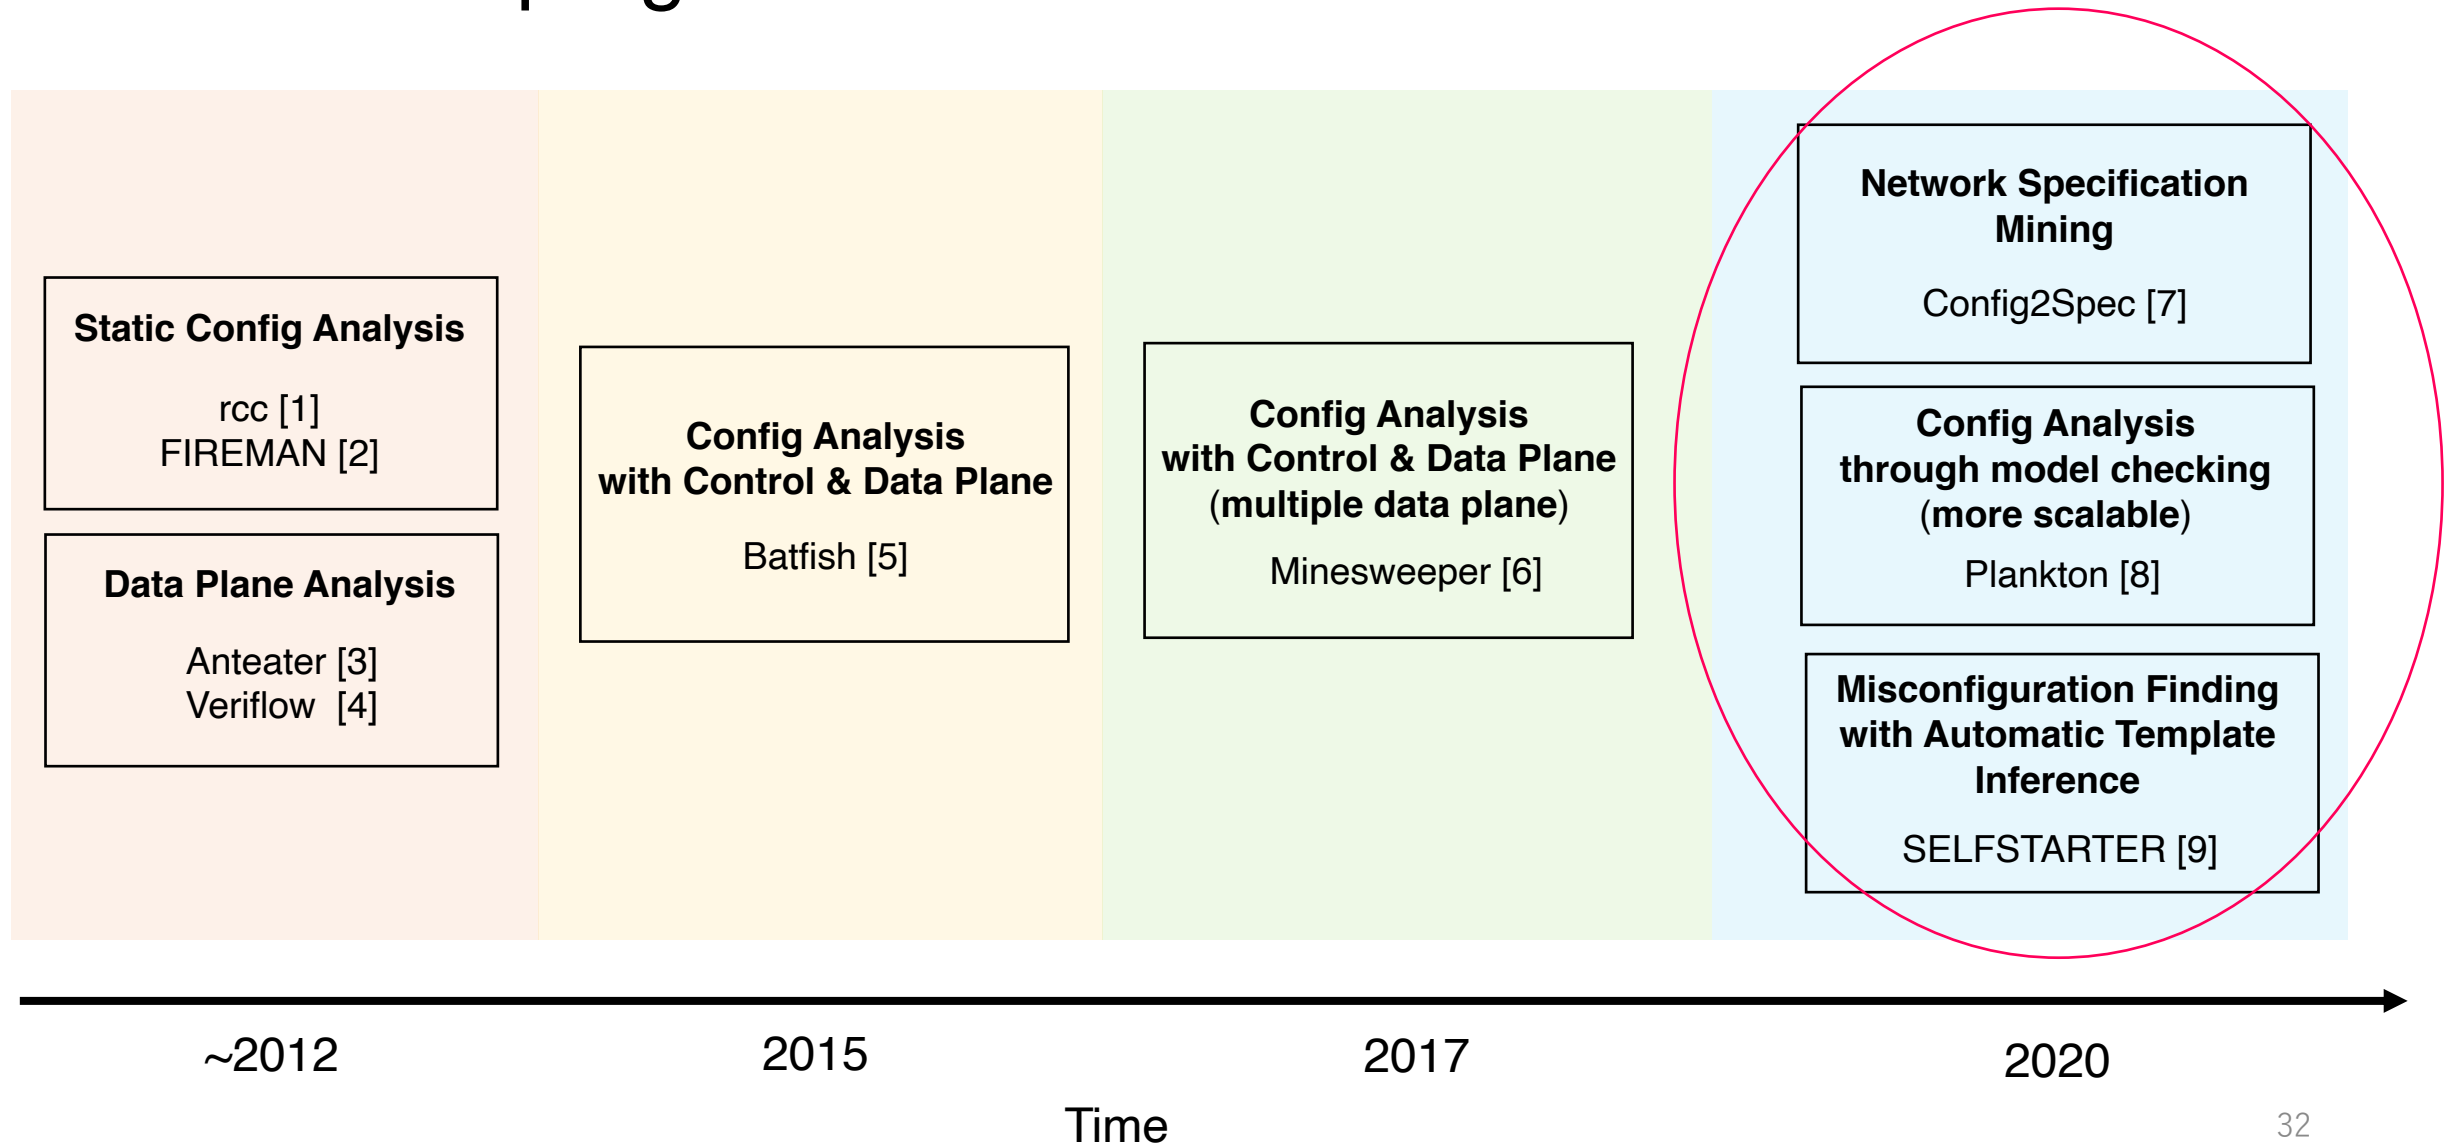
**Black holes only occur
at the network edge**

less than 1.5sec

Summary



Historical progress



Short summary of the three works

Plankton [8]

- A verification tool to scale better than the SMT solver approach utilizing a model checker and domain-specific optimizations

Config2Spec [7]

- Automatically synthesizing a formal specification of a network given its configuration and a failure model

Short summary of the other works

SELFSTARTER [9]

- Automatically identifying configuration outliers by template inference from the configuration

Discussion

- Research such as knowledge extraction from configuration has been popular (not just about scalability and coverage)
- So many research focuses on the reachability
- IPv6 ...?

My future Work

- Read the three papers
- Skim through previous papers that have not been yet read

Reference

- [1] Detecting BGP Configuration Faults with Static Analysis
Nick Feamster and Hari Balakrishna (NSDI'05)
- [2] FIREMAN: A Toolkit for FIREwall Modeling and Analysis
Lihua Yuan et al., (S&P'06)
- [3] Debugging the Data Plane with Anteater
Haohui Mai et al., (Sigcomm'2011)

Reference

[4] Veriflow : Verifying Network-Wide Invariants in Real Time

Ahmed Khurshid et al.,(NSDI'13)

[5] A General Approach to Network Configuration Analysis

Ari Fogel et al., (NSDI'15)

[6] A General Approach to Network Configuration Verification

Ryan Beckett et al., (Sigcomm'2017)

Reference

- [7] Config2Spec: Mining Network Specifications from Network Configurations Rüdiger Birkner et al.,(NSDI'20)
- [8] Plankton: Scalable network configuration verification through model checking, Santhosh Prabhu et al., (NSDI'20)
- [9] Finding Network Misconfigurations by Automatic Template Inference, Siva Kesava Reddy Kakarla et al., (NSDI'20)